CSE 321b

Computer Organization (2)

تنظيم الحاسب (2)

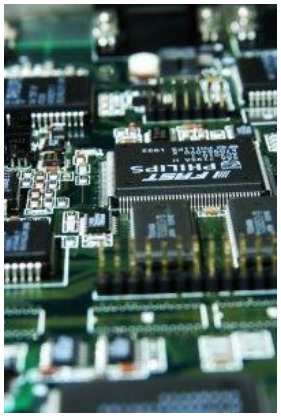

3rd year, Computer Engineering

Winter 2017

Lecture #2

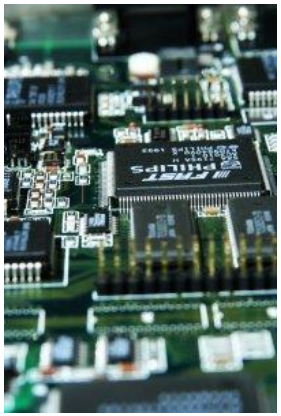

Dr. Hazem Ibrahim Shehata

Dept. of Computer & Systems Engineering

Credits to Dr. Ahmed Abdul-Monem Ahmed for the slides

Adminstrivia

- Schedule:
  - —Lectures: Wednesday 10:15am – 12:45pm
  - —Tutorials: after lecture (this week only)
  - —Office hour: TBA
- Assignment #1:
  - —To be released next week

Website: http://hshehata.github.io/courses/zu/cse321b
Office hours: TBA

Ch 5: Internal Memory (*Cont.*)

Error Correction

- Semiconductor memory is subject to errors
  - Rate: 1 error/hour to 1 error/century in a 1GB memory!
  - Types: hard and soft.
- Hard Failure
  - Permanent physical defect.
  - Mem. cells can't store data: stuck at 0 or 1, or switching.
  - Caused by harsh environments, manufacturing defects, or wear.
- Soft Error
  - Random, non-destructive event that alters contents of one or more memory cells.
  - No permanent damage to memory.
  - Caused by power supply problems or alpha particles.
- Detected/corrected using Hamming error correcting code.

Error-Correcting Code Function

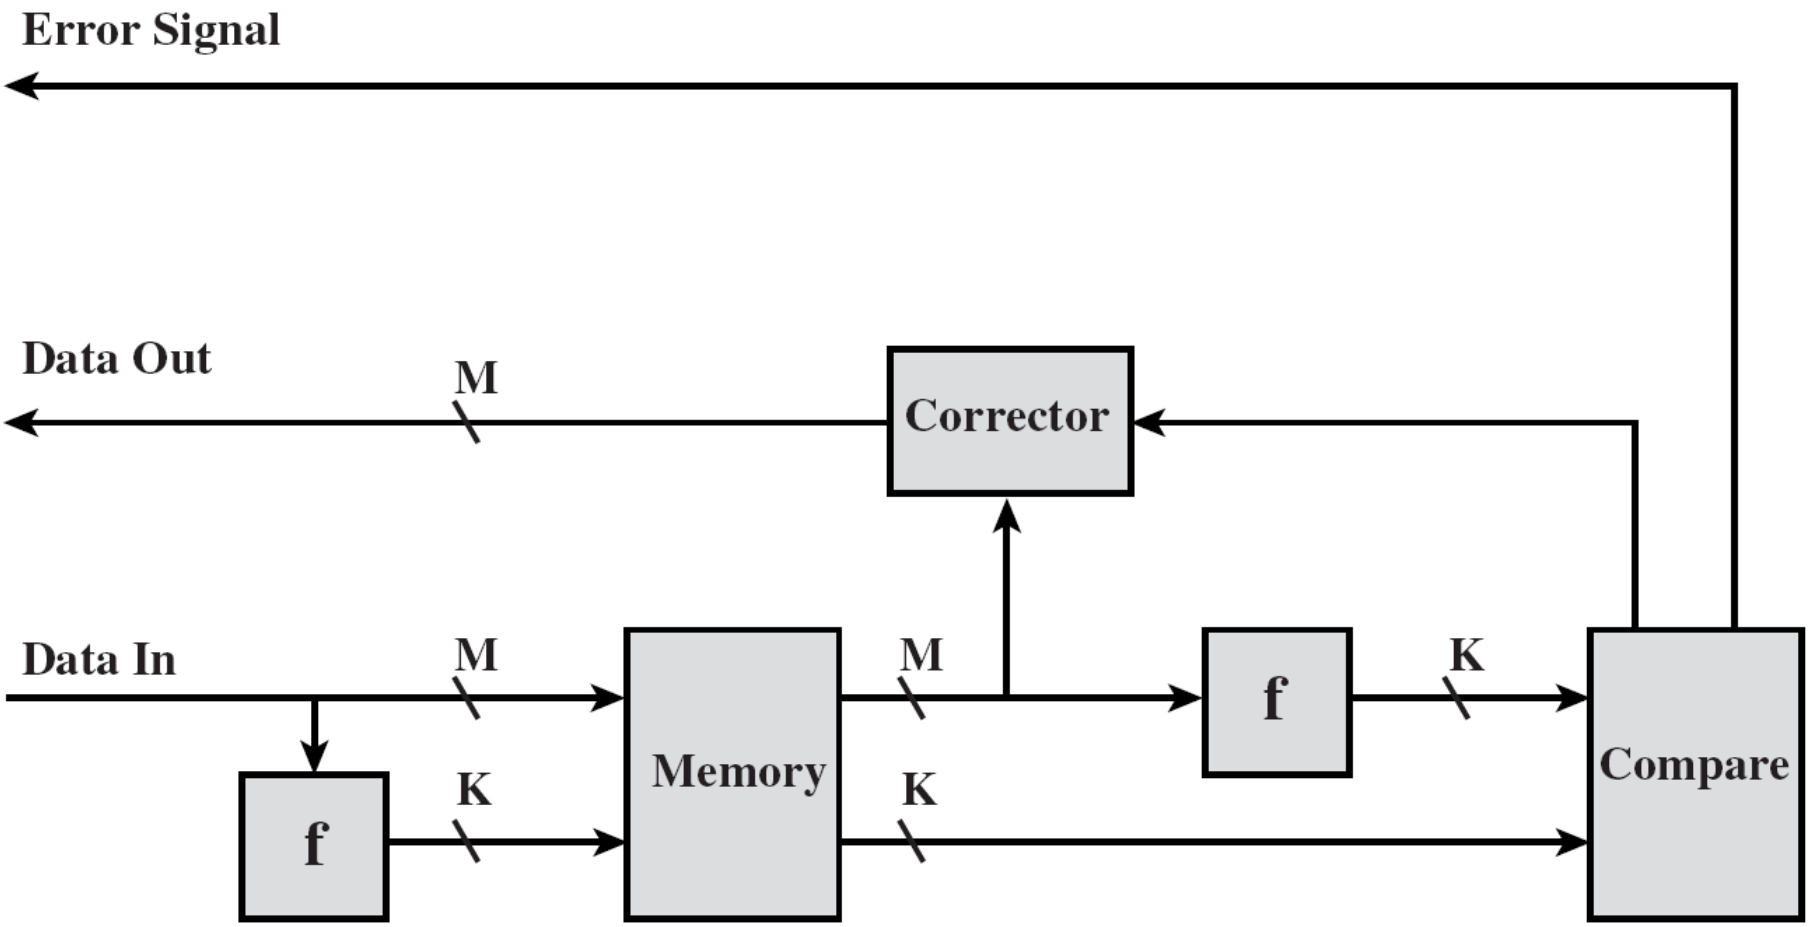

Hamming Error-Correcting Code

Data bits: 1110

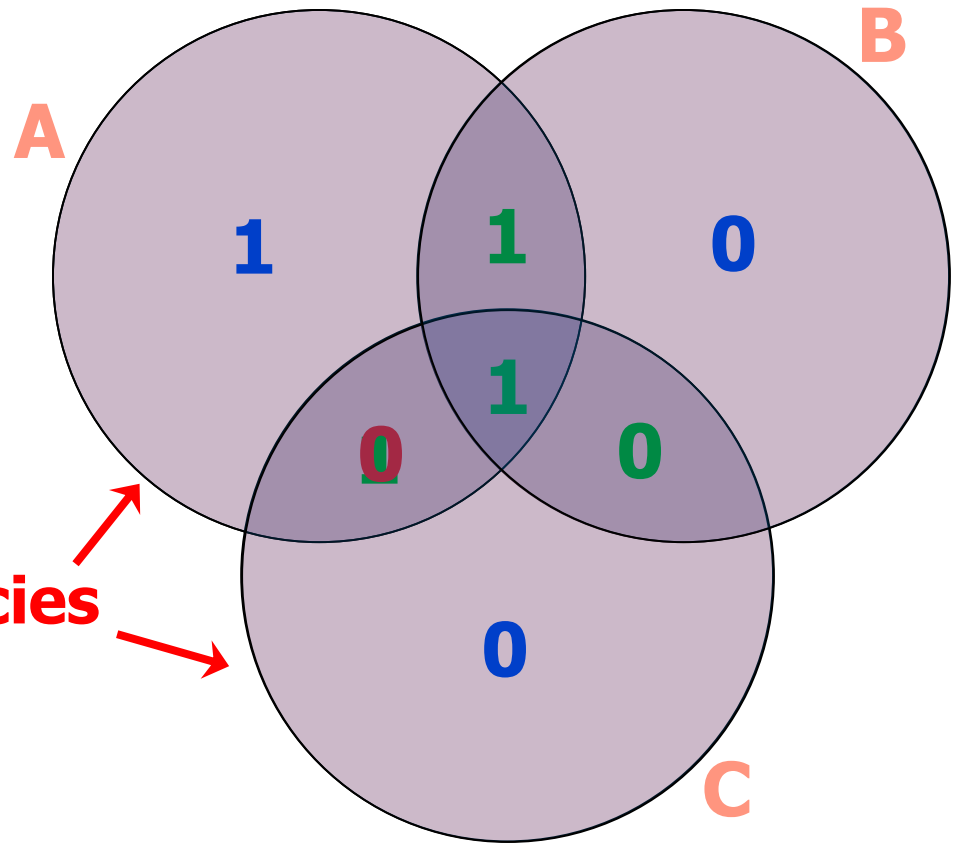

Discrepancies

Parity bits

Chosen so that total number of 1s in each circle is even.

By checking the parity bits, discrepancies are found ➔ error can be easily found and corrected.

Error-Correcting Codes

- A codeword consists of N bits split into M data bits and K check (redundant) bits
  - $N = M + K$.
- Hamming distance: Number of bit-positions in which two codewords differ.
  - Ex.: 11001001, 10100001 → Hamming distance = 3 → 3 bit errors are needed to convert one into the other.
- Note: in a code, not all 2^N codewords are legal.
- Hamming distance of the whole code: minimum Hamming distance between 2 legal codewords.
- A distance d code can:
  - **Detect:** $d - 1$ errors.
  - **Correct:** $(d - 1)/2$ errors if d is odd, or $(d/2) - 1$ errors if d is even.

Error Detection/Correction

- Detection: parity bit.
  - Distance = 2 → can detect up to 1 bit error.
  - Ex.: data=1011010 → codeword=10110100 → any codeword with a distance = 1 (such as: 101**0**0100) is considered illegal → single-bit errors are detectable.
- Correction: Consider a code with 4 valid codewords:

  0000000000, 0000011111, 1111100000, 1111111111
  - Distance = 5 → can correct up to 2 bit errors.
  - If 0000000111 arrives, → 0000011111
  - If 0000000000 becomes 0000000111 due to 3 errors → cannot be corrected properly.

Single Bit Error Correction

Design a code to correct all single bit errors.

- M = data bits, K = check bits
- $N = M + K$
- **Rule**: Choose K s.t. $\boldsymbol{M + K + 1 \leq 2^K}$
- **Justification**:
  - Each of the 2^M legal words has N illegal codewords at distance 1.
  - Thus, each of the 2^M legal words requires $(N + 1)$ bit patterns dedicated to it.
  - $(N + 1)\ 2^M \leq 2^N \rightarrow M + K + 1 \leq 2^K$

Hamming Code

$2^K - 1 \geq M + K$ ➔ $2^K \geq 9+K$ ➔ $K = 4$

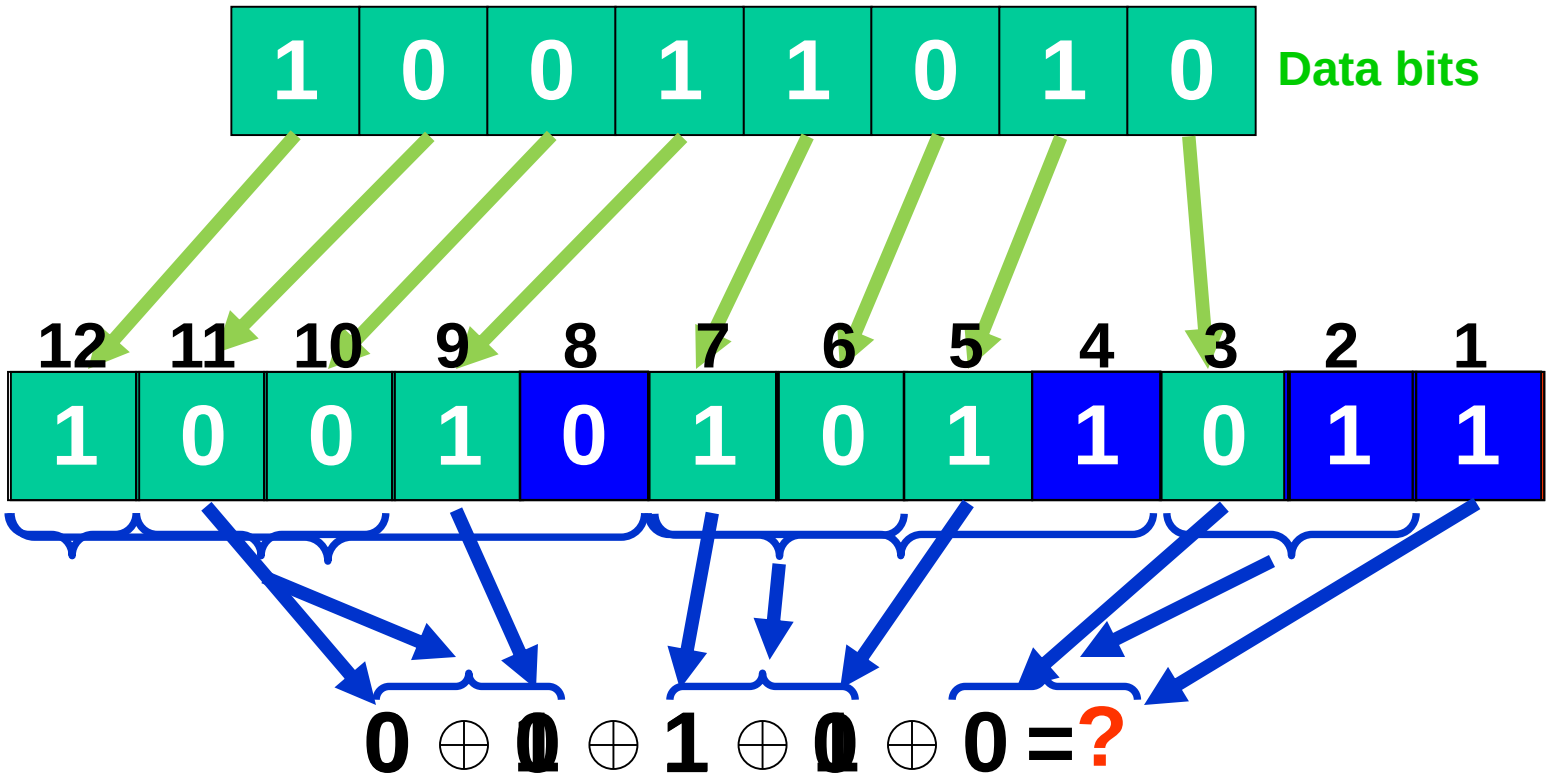

Bit position 8: 0

Hamming Code (2)

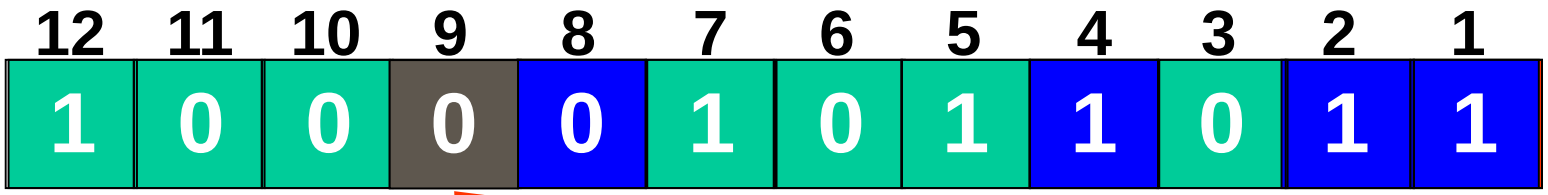

- Assume error in bit 9.
- Recompute the check bits.
- Bit 1 = 0 (error).
- Bit 2 = 1.
- Bit 4 = 1.
- Bit 8 = 1 (error).
- Error is in bit position = 1 + 8 = 9 → flip it (correction).

Hamming SEC-DED Code

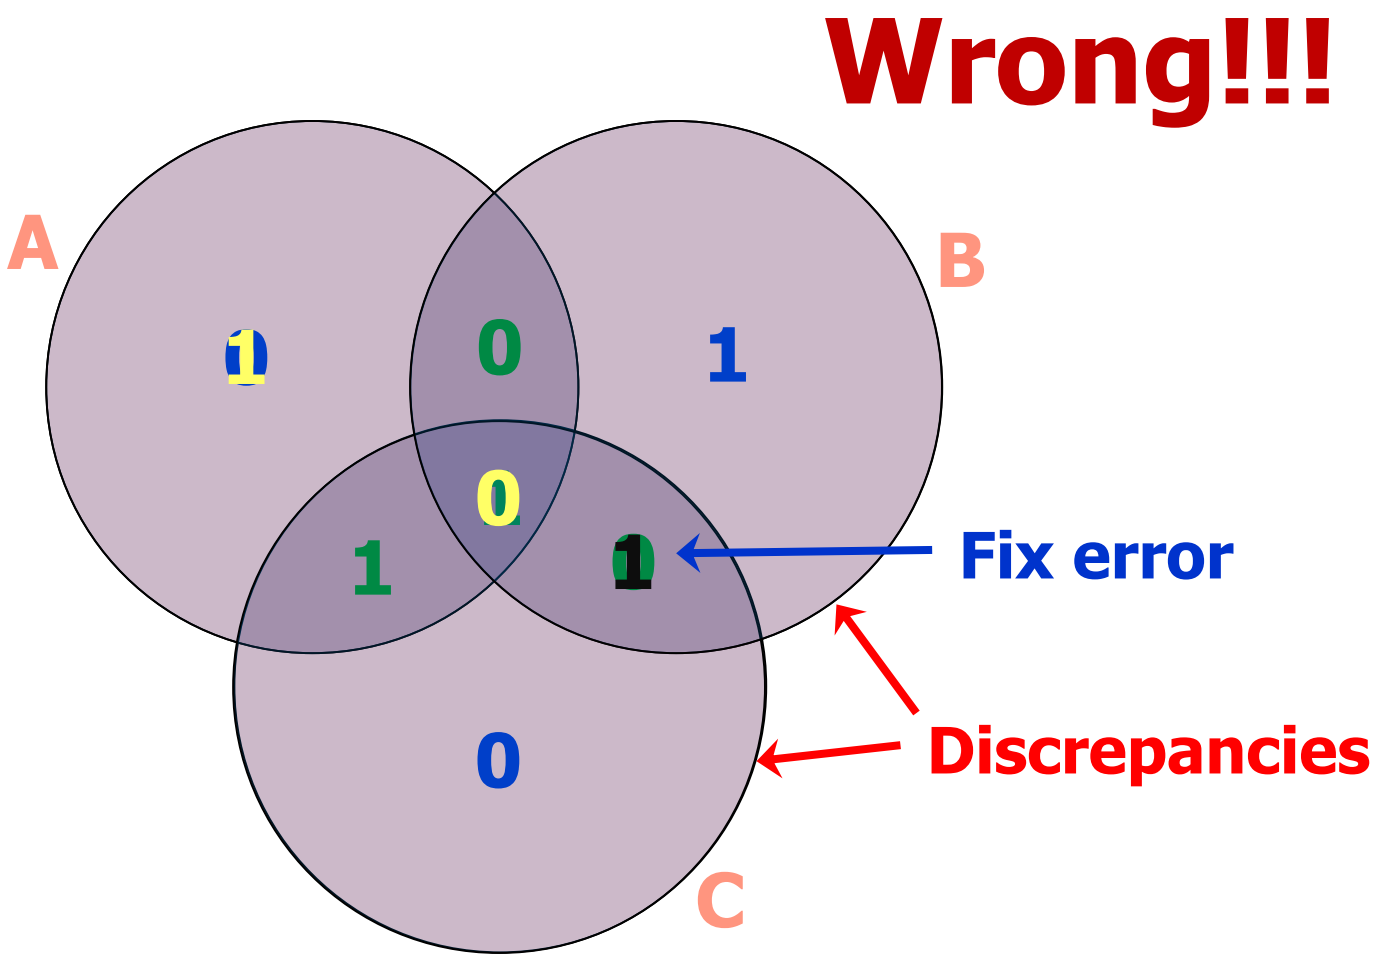

Hamming SEC-DED Code (2)

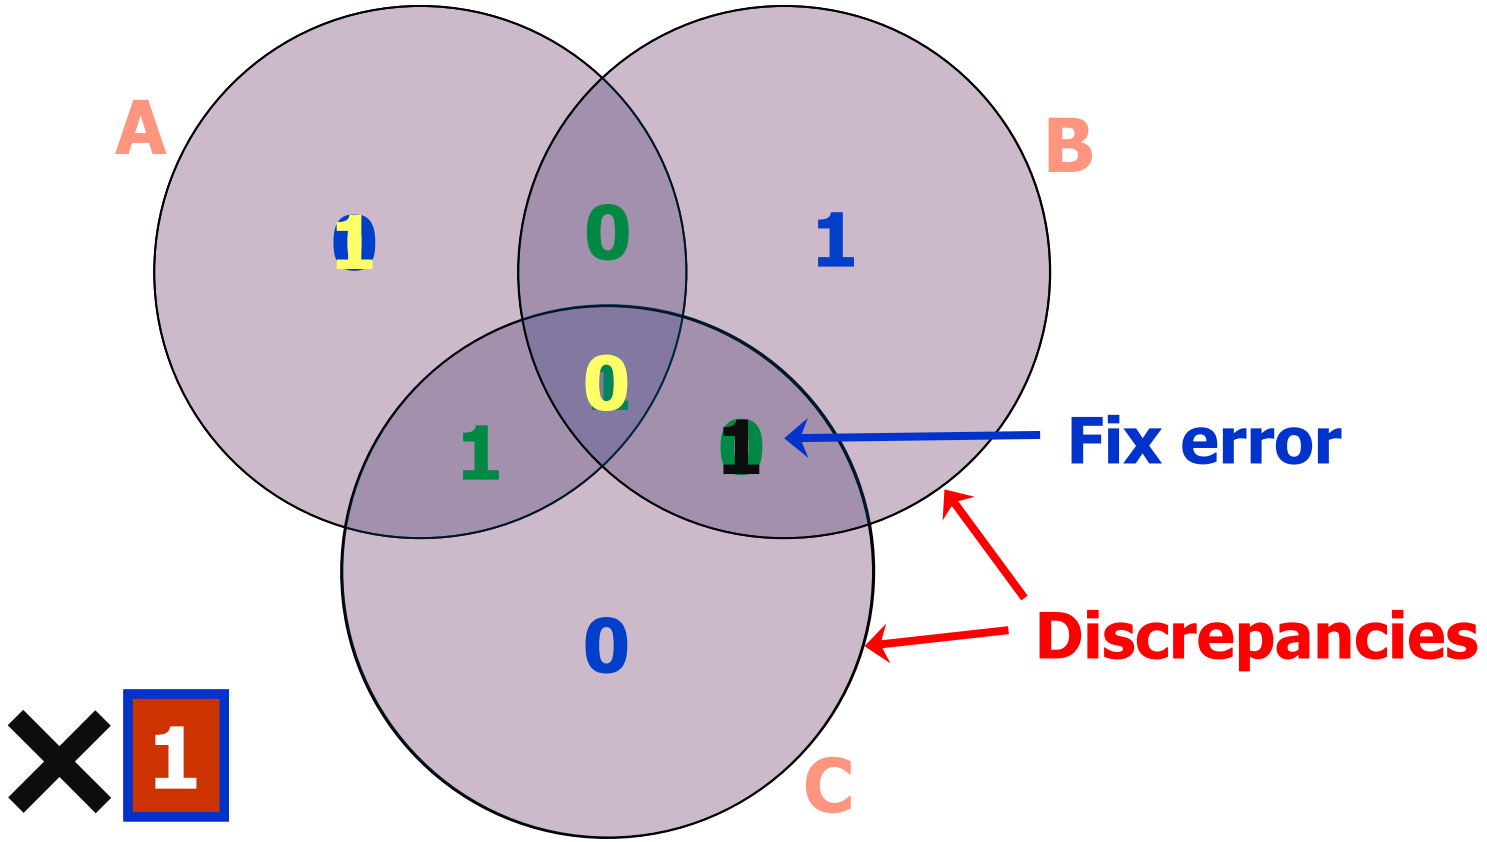

Increase in Word Length with Error Correction

	Single-Error Correction		Single-Error Correction/ Double-Error Detection	
Data Bits	**Check Bits**	**% Increase**	**Check Bits**	**% Increase**
8	4	50	5	62.5
16	5	31.25	6	37.5
32	6	18.75	7	21.875
64	7	10.94	8	12.5
128	8	6.25	9	7.03
256	9	3.52	10	3.91

Advanced DRAM Organization

- Interface to MM is a system bottleneck.
- DRAM chip is the main building block of MM.
- Basic DRAM architecture same since 1970s!
- Enhancements to basic DRAM architecture
  - —Synchronous DRAM (SDRAM)
    - DDR-SDRAM, DDR2-SDRAM, DDR3-SDRAM
  - —Rambus DRAM (RDRAM)
  - —Cache DRAM (CDRAM)

Synchronous DRAM (SDRAM)

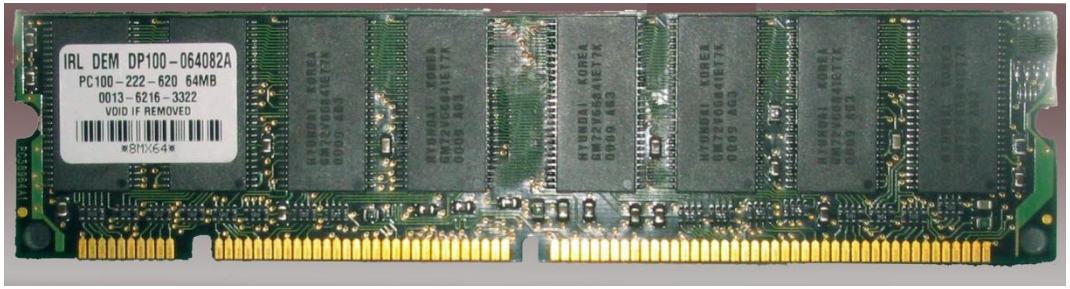

- Unlike traditional DRAM (which is asynchronous), SDRAM exchanges data with CPU synchronized to an external clock (system bus).
- No wait states!
  - —SDRAM moves data in/out under control of system clock.
  - —CPU issues command and address.
  - —Latched by the SDRAM.
  - —SDRAM responds after a number of clock cycles.
  - —Meanwhile, CPU can do other tasks ➔ no waits.

Synchronous DRAM (SDRAM) (2)

- Burst mode: a series of data bits can be clocked out rapidly after the first bit has been accessed.
- Eliminates row and column address setup time.
- Useful when the required bits are in sequence and in the same row of the array as the initial access.
- Multiple-bank internal architecture improves opportunities for parallelism.
- Mode Register (MR)
  - Specifies burst length.
  - Allows programmer to adjust latency between read request and data transfer.

IBM 64Mb SDRAM

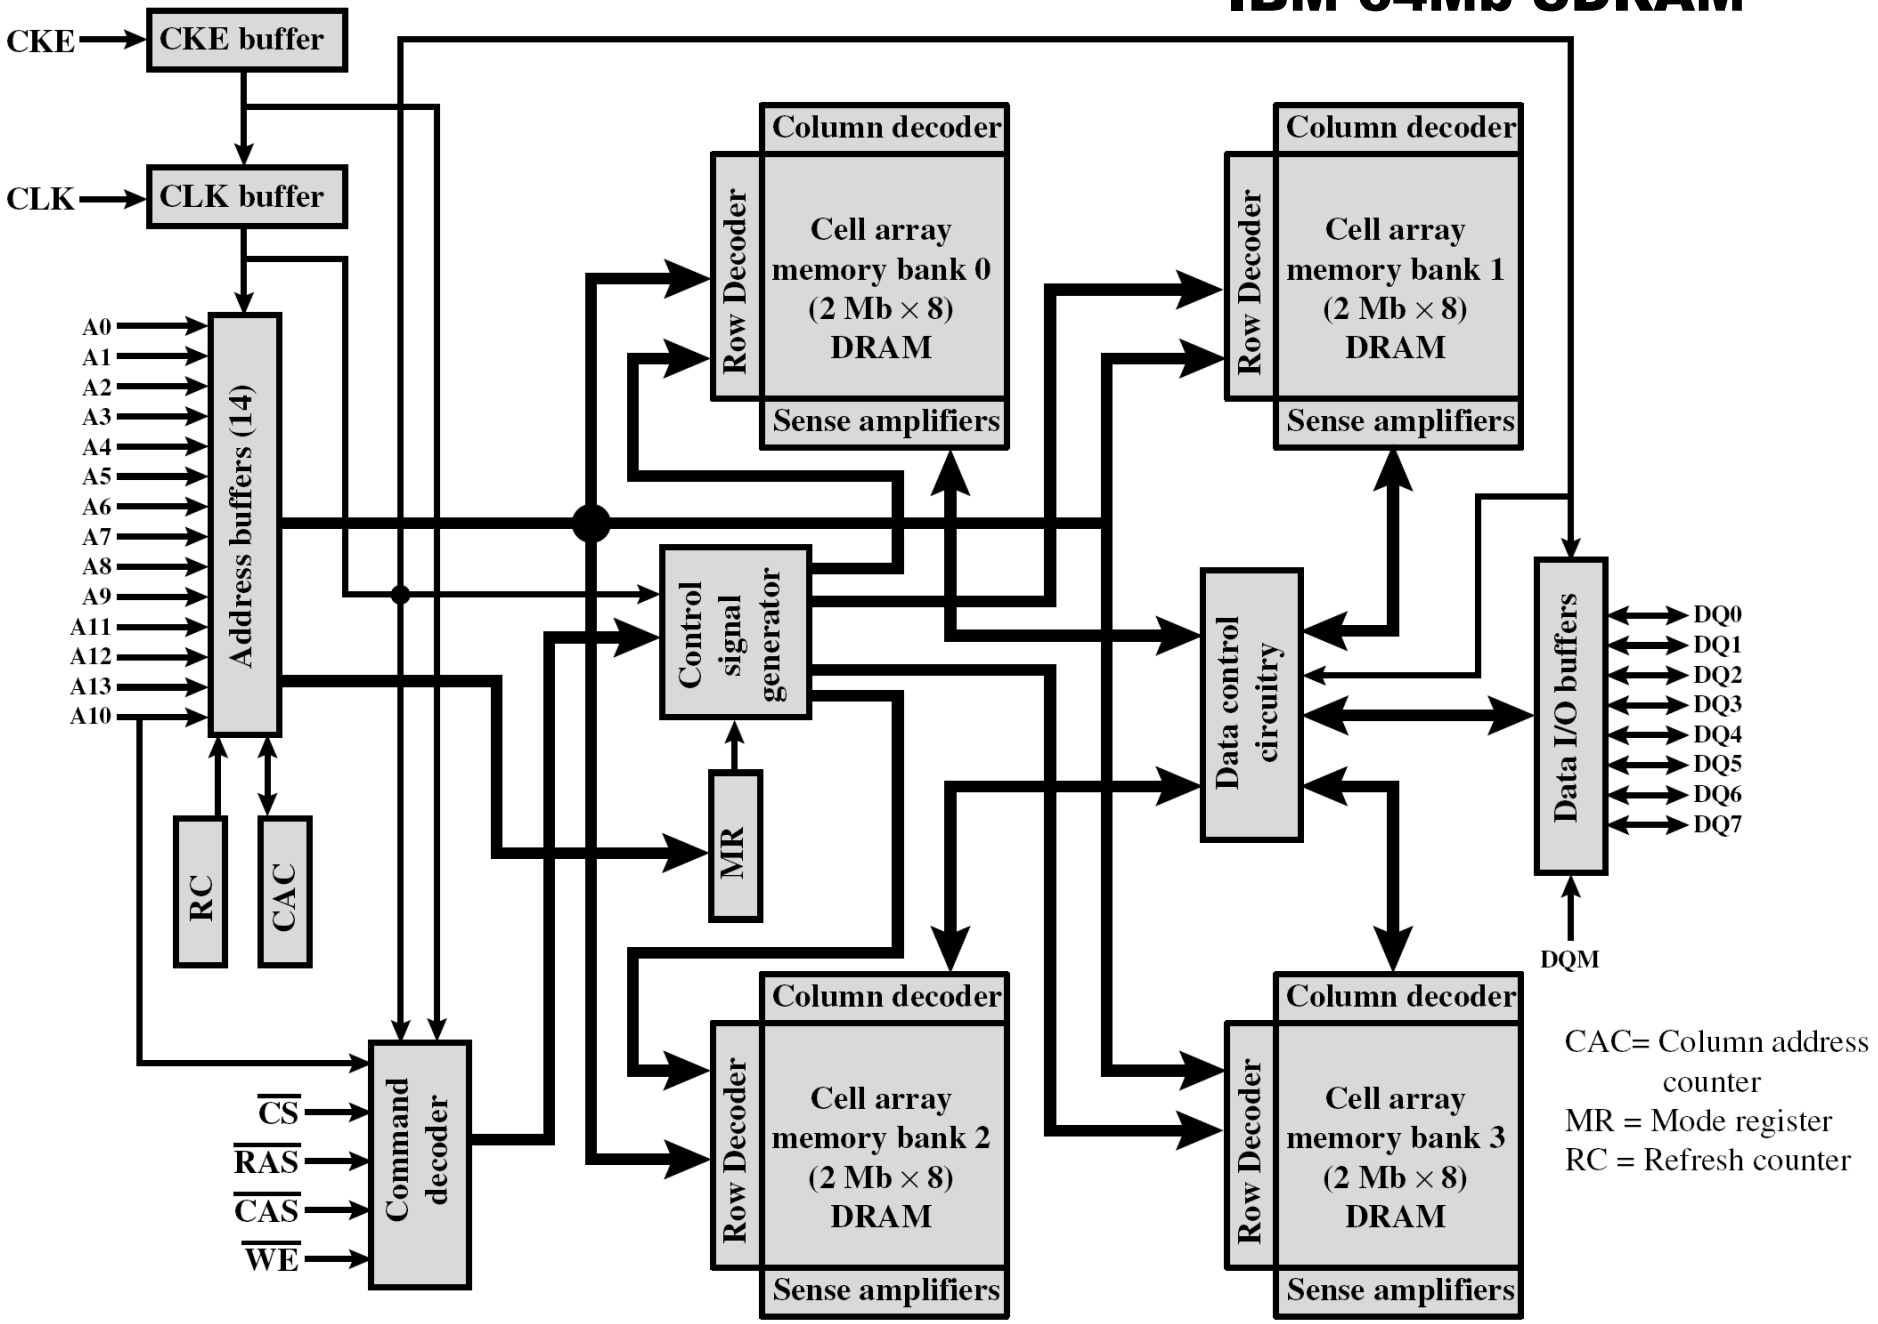

SDRAM Operation

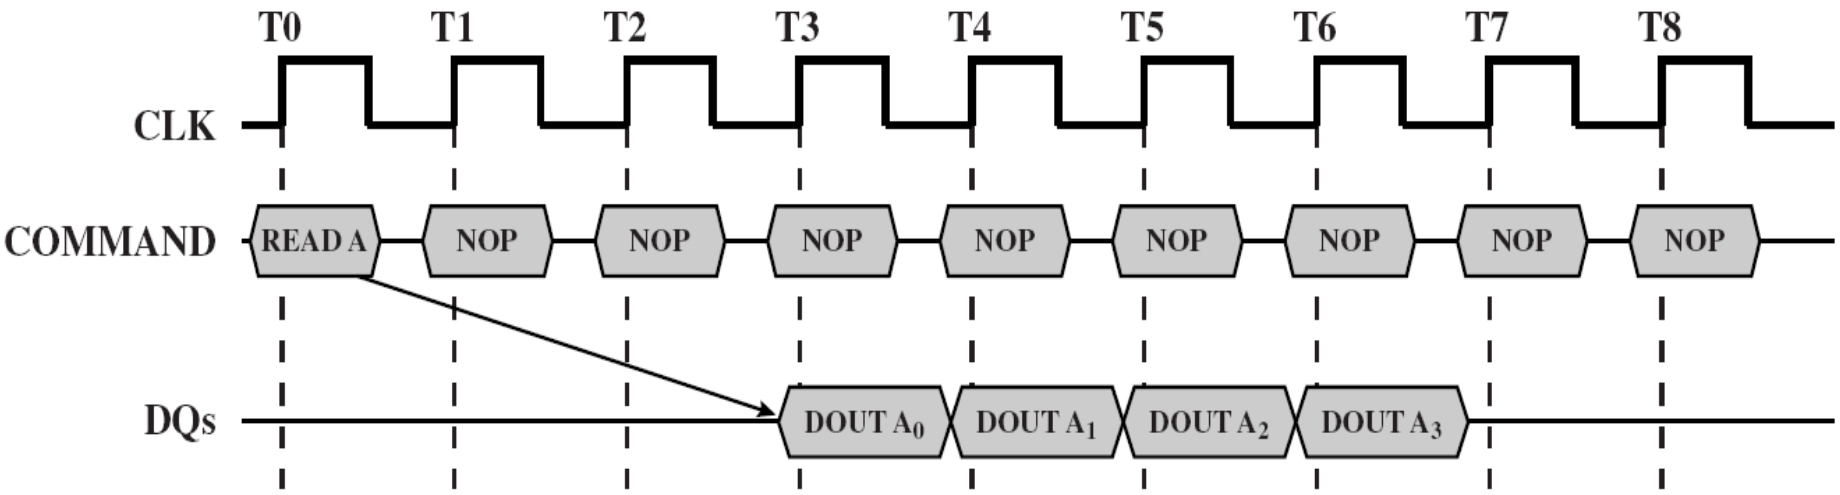

- Burst type, length and latency are set in mode reg.
  - —Type: interleaved or sequential
  - —In the example: length=4 and latency=3
- Burst read command is initiated.
  - —At the rising edge of the clock: CS & CAS ➔ low, and RAS & WE ➔ high.
- Address inputs determine starting column for burst.

Enhanced versions of SDRAM

- Double Data Rate (DDR-SDRAM)
  - —produces two words of data every memory cycle
  - —sends data to CPU twice per clock cycle (at rising & falling edges of the clock)
  - —2x data rate of SDRAM (with same cell speed)
- DDR2-SDRAM
  - —produces four words of data every memory cycle
  - —2x bus speed & data rate of DDR (with same cell speed)
- DDR3-SDRAM
  - —produces eight words of data every memory cycle
  - —2x bus speed & data rate of DDR2 (with same cell speed)

Reading

Rambus DRAM (RDRAM) (1)

- Adopted by Intel for Pentium & Itanium in late 90s.
- Was main competitor to SDRAM.
- Chips are vertical packages – all pins on one side.
- Data exchange over 28 wires < 12 cm long.
- Bus can address up to 32 RDRAM chips and is rated at 1.6 GBps.

Reading

Rambus DRAM (RDRAM) (2)

- Asynchronous block-oriented protocol
  - Initial 480 ns access time.
  - Then 1.6 GB/s.
- What makes this speed possible is the bus itself, which defines impedances, clocking, and signals very precisely.
- RDRAM gets a memory request over the high-speed bus (unlike conventional DRAMs controlled by RAS, CAS, R/W, and CE).
- This request contains address, operation, number of bytes.

Reading

RDRAM Diagram

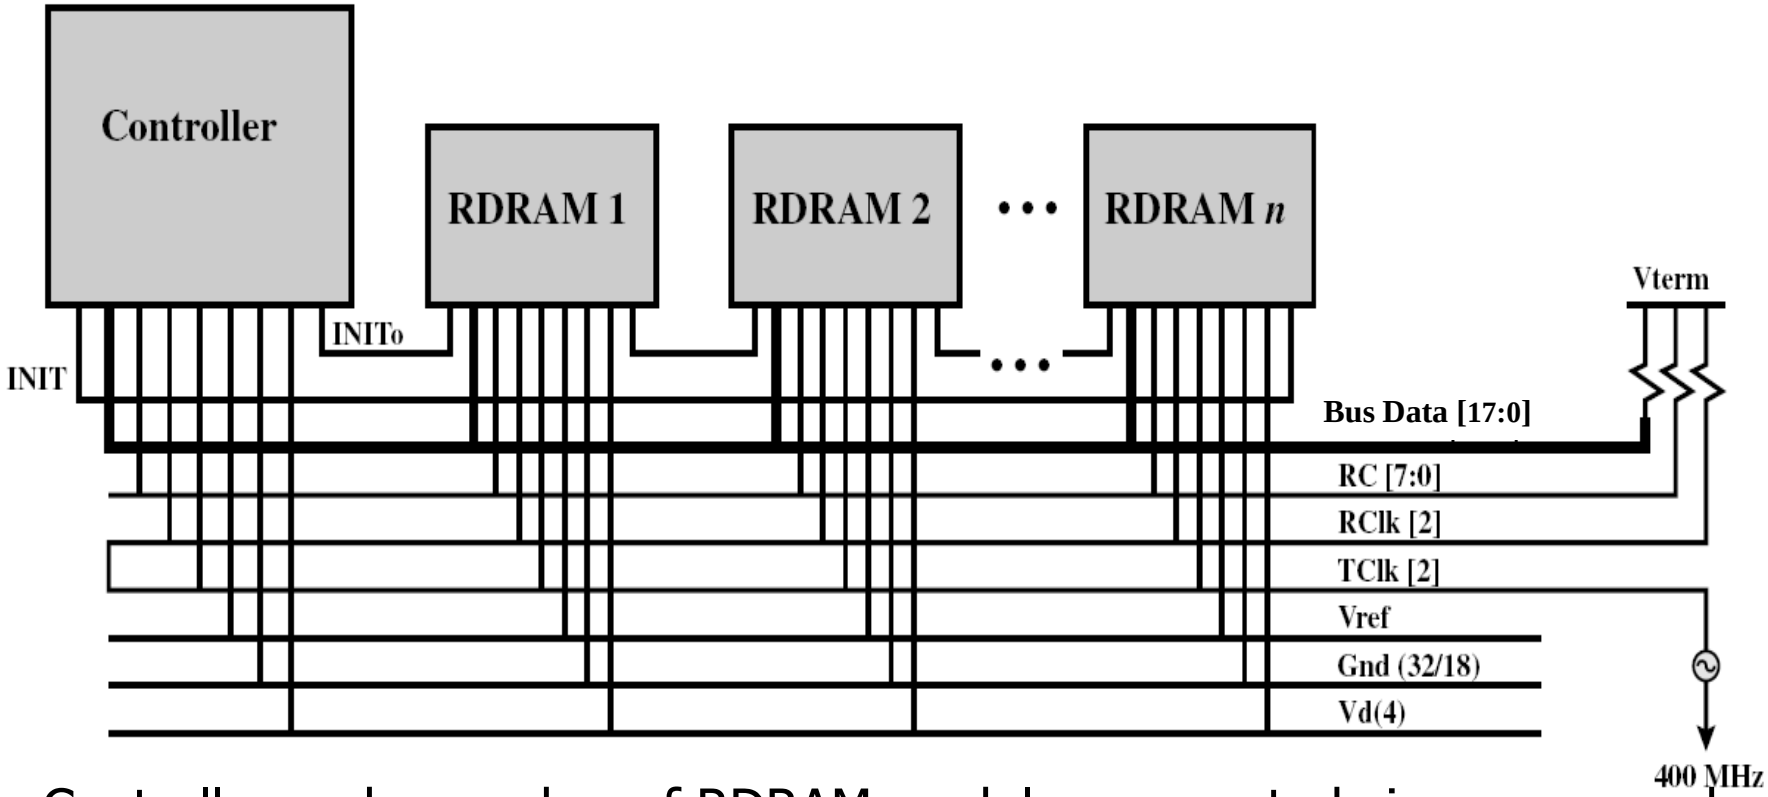

- Controller and a number of RDRAM modules connected via a common bus.
- Bus: 18 data lines, cycling at twice the clock rate ➔ 800 Mbps per line.
- Address and control signals: 8 lines (RC)
- Clock starts at the far end from the controller, propagates to the controller end, then loops back.
- Module sends data to the controller synchronously to the clock to master.
- Controller sends data to a module synchronously with the clock in the opposite direction.

Reading

Cache DRAM (CDRAM)

- CDRAM integrates a small SRAM cache (16 kb) onto a generic DRAM chip.
- CDRAM can be used in two ways:
  1. As a **true cache** with 64-bit lines.
  2. As a **buffer** to support serial access of a data block.
     - e. g., to refresh a bit-mapped screen, the CDRAM can prefetch the data from the DRAM into the SRAM buffer. Subsequent accesses to the chip results in accesses to the SRAM only.

Reading Material

- Stallings, Chapter 5:
  - Pages 170 – 180