

CSE 321b


Computer Organization (2)

تنظيم الحاسب (2)



3rd year, Computer Engineering
Winter 2016

Lecture #7



Dr. Hazem Ibrahim Shehata

Dept. of Computer & Systems Engineering

Credits to Dr. Ahmed Abdul-Monem Ahmed for the slides

Adminstrivia

- Assignment #2:
 - To be released later on this weekend.
- Midterm:
 - Date: Thursday, April 21, 2014
 - Time: 10:30am – 12:00pm
 - Location: classrooms #27321 (قاعة 4) & #27309
 - Coverage: lectures #1 → #7

Website: <http://hshehata.github.io/courses/zu/cse321b/>

Office hours: Sunday 11:30am – 12:30pm

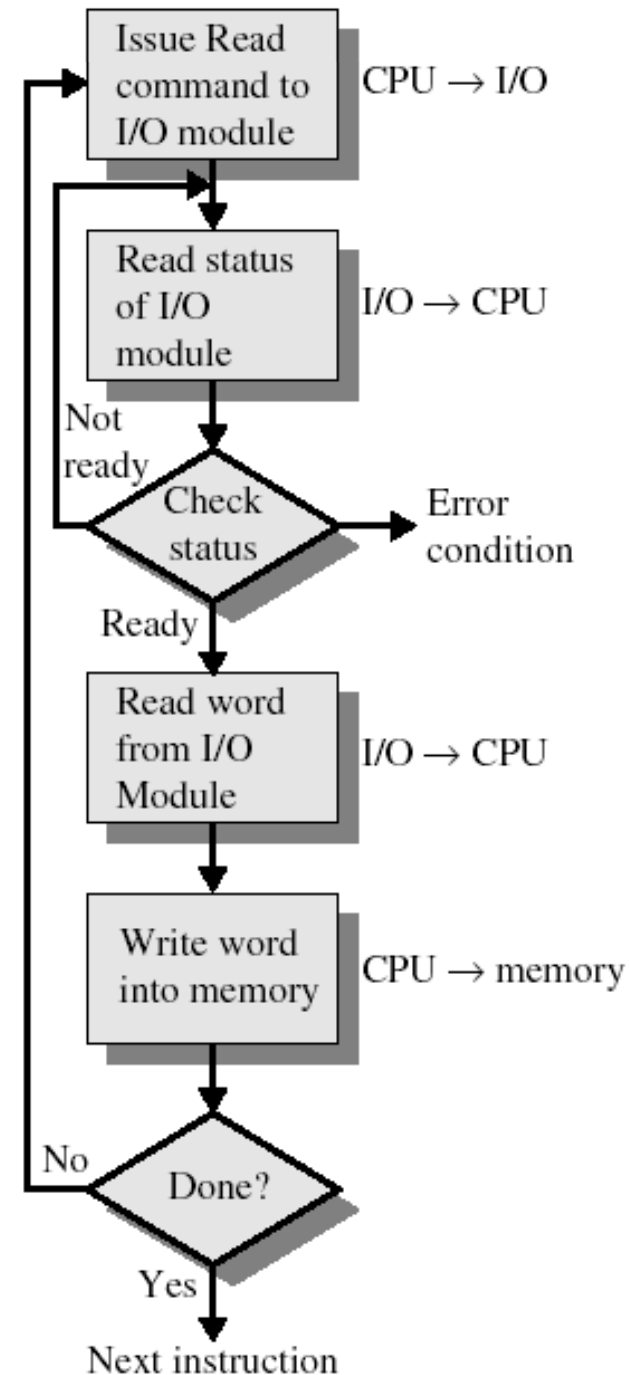
Chapter 7. Input / Output (*Cont.*)

Outline

- External Devices
 - Types
 - Structure
- I/O Modules
 - Function
 - Structure
- I/O Techniques
 - Programmed I/O
 - Interrupt-Driven I/O
 - Direct Memory Access
- I/O Channels & Processors

Programmed I/O

- CPU **issues** read/write command to I/O module.
- CPU **periodically checks** status bits.
 - CPU **waits** for command to complete.
 - I/O module **does not interrupt** CPU.
- CPU reads data from device/memory and writes it to memory/device.

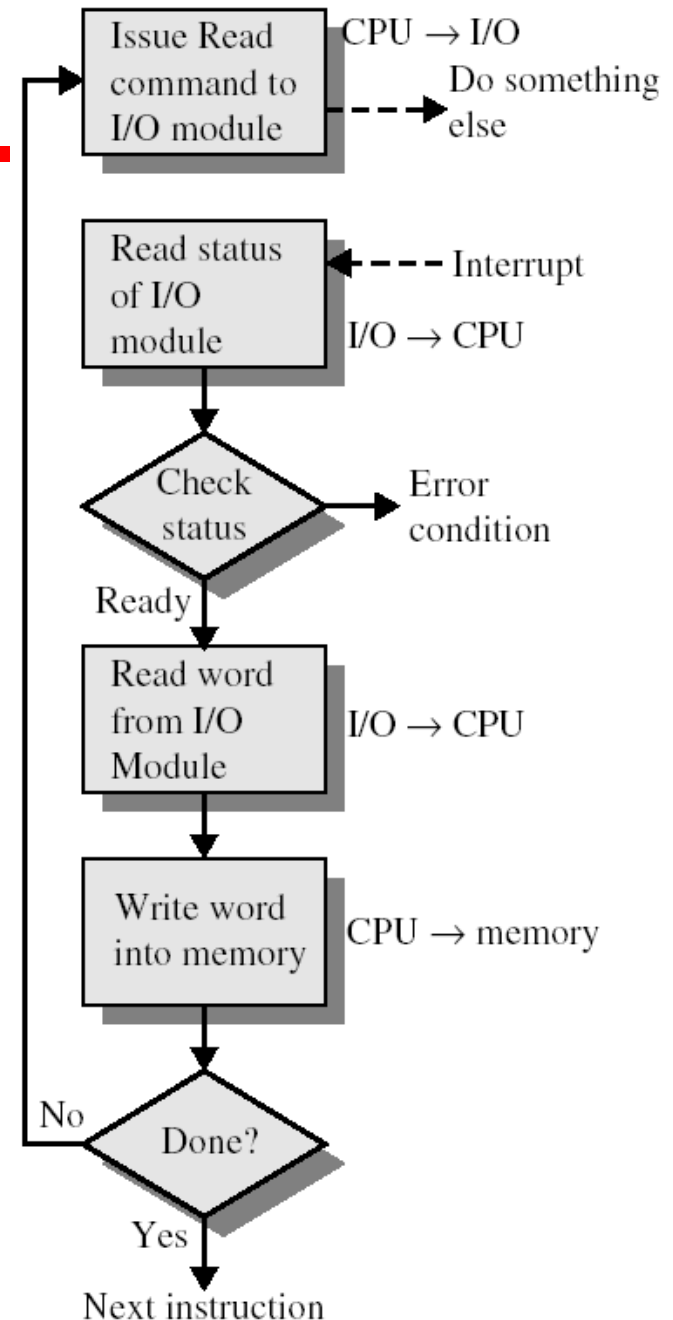


Interrupt-Driven I/O

- **Purpose:** To overcome CPU waiting.
- No repeated CPU checking of device.
- CPU issues command and moves on to do other useful work.
- I/O module **interrupts** CPU when ready.

Interrupt-Driven I/O

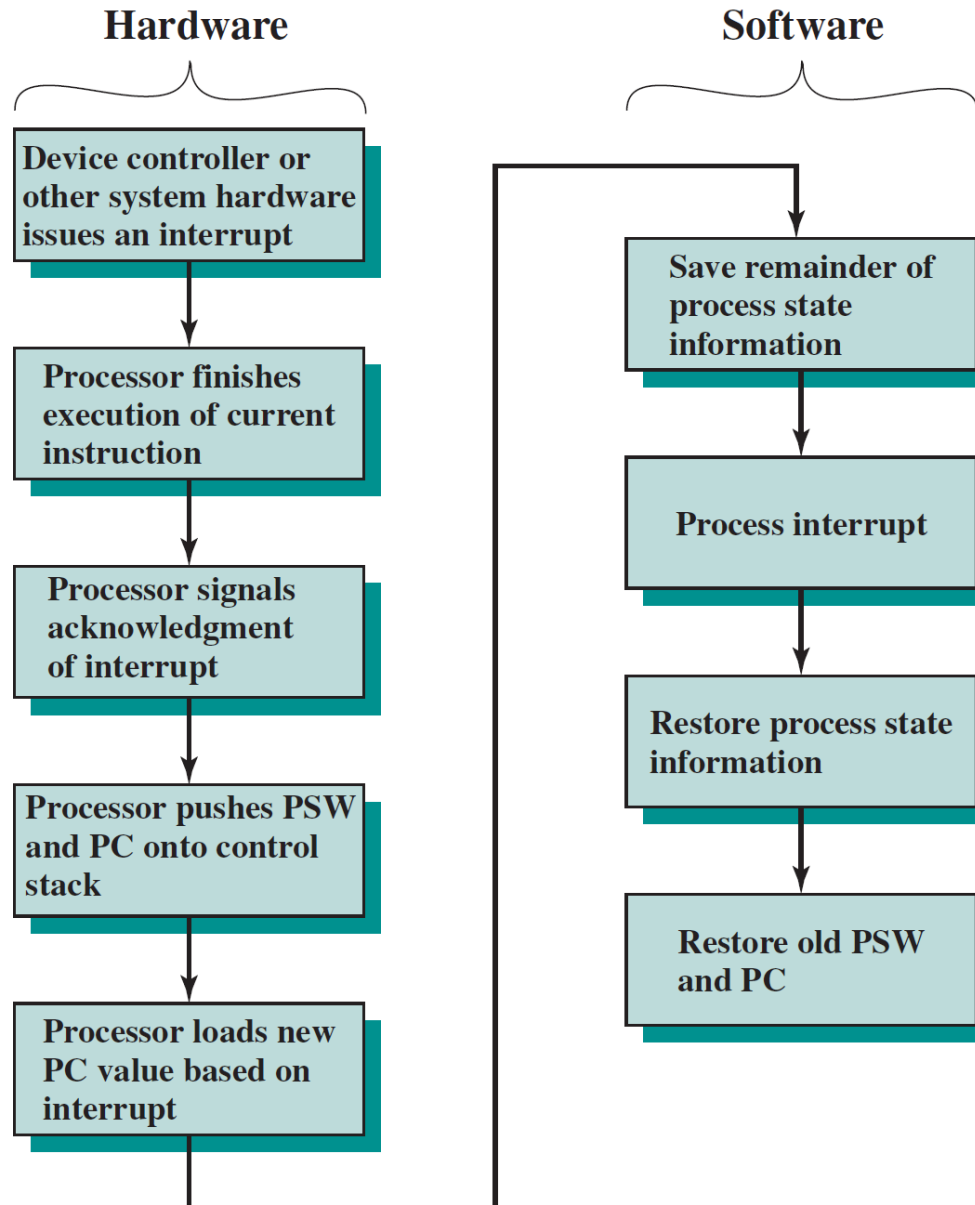
- CPU **issues** read command
- I/O module gets data from peripheral whilst CPU **does other work**.
- I/O module **interrupts** CPU.
- CPU **reads** data from I/O module.
- CPU **writes** data to memory.



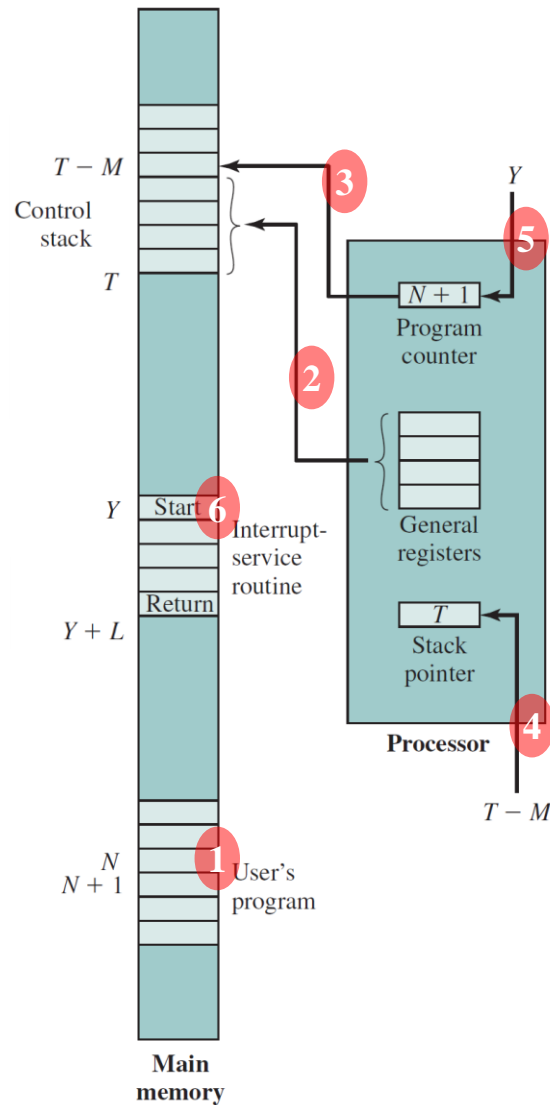
CPU Viewpoint

- Issue read command.
- Do other work.
- Check for interrupt at end of each instruction cycle.
- If interrupted:
 - Save context (registers).
 - Process interrupt.
 - Fetch data (from module) & store (to memory)

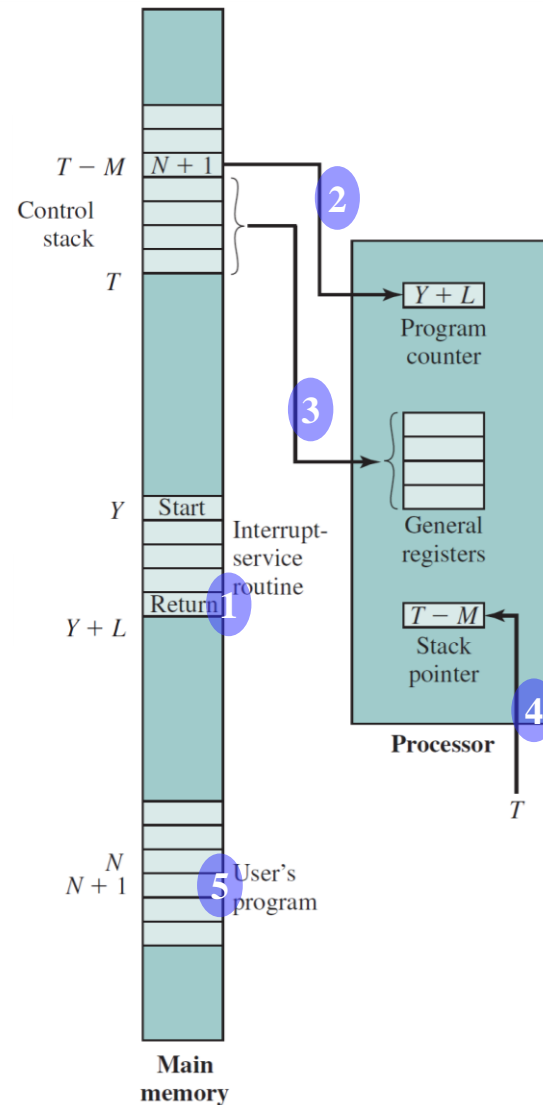
Simple Interrupt Processing



Changes in Mem. & Reg.'s upon Interrupt



(a) Interrupt occurs after instruction at location N



(b) Return from interrupt

Identifying Interrupting Module

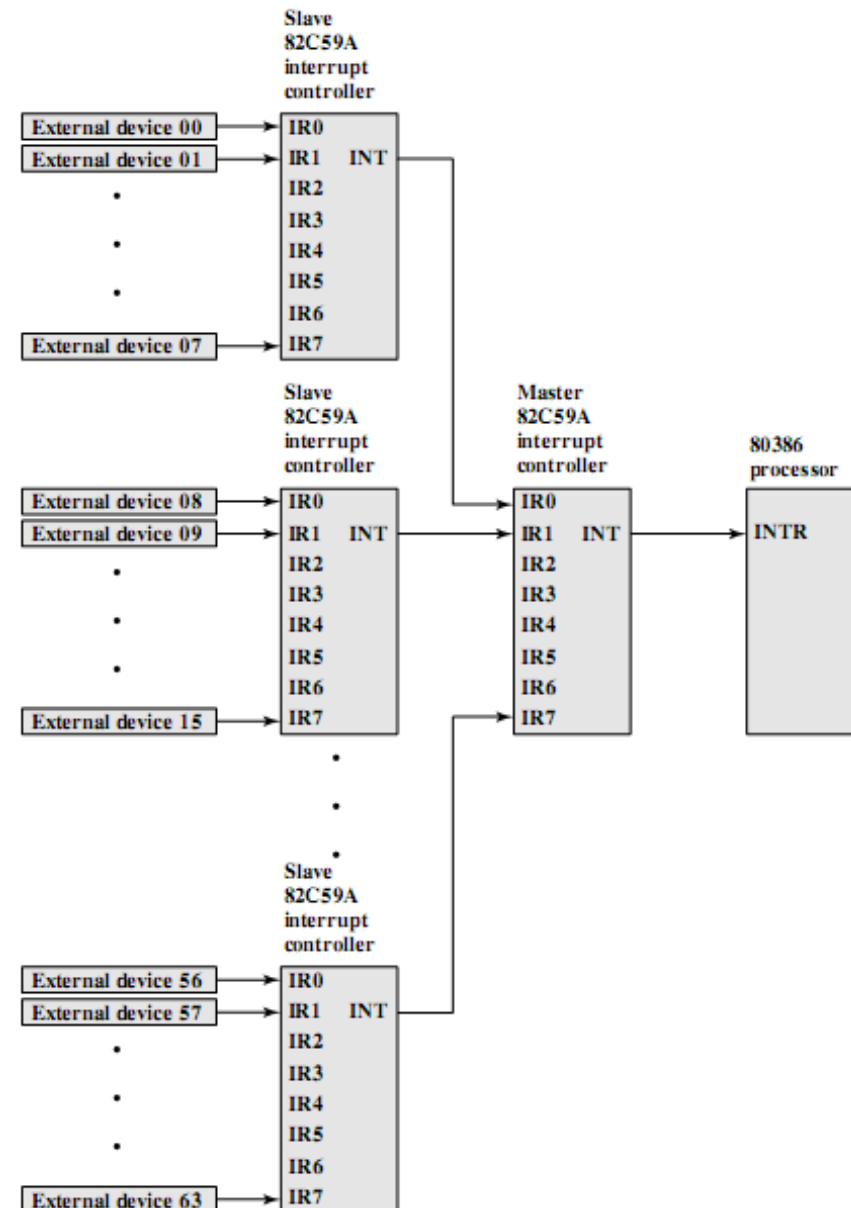
- How to identify the module issuing the interrupt?
 1. Different line for each module
 - Limits number of devices.
 2. Software poll (single line)
 - CPU asks each module in turn → time consuming!!
 - e.g., Send TESTI/O → Set address lines → Check status reg.
 3. Daisy chain or Hardware poll (single line)
 - All modules share a single interrupt request line.
 - Interrupt acknowledge sent down a chain.
 - Module (that issued interrupt signal) places its vector on bus.
 - CPU uses vector to identify handler routine.
 4. Bus Master (single line)
 - Module must claim the bus before it can raise interrupt.
 - e.g., PCI & SCSI.

Multiple Interrupts & Priorities

- How to deal with simultaneous interrupts?
- Solution: prioritize interrupts!
 1. With multiple lines:
 - Each interrupt line has a priority.
 - Higher priority lines can interrupt lower priority lines.
 2. With software polling:
 - Priority determined by order in which modules are polled.
 3. With daisy chain:
 - Priority determined by order of modules on chain.
 - Closer modules have higher priority.
 4. With bus arbitration:
 - Only current master can interrupt.
 - Priority is defined by the bus arbitration protocol.

Example - PC Bus

- 80386 has one interrupt line!
- To handle more interrupts, connect 1 (or more) interrupt arbiter → **Intel 8259**.
- Intel 8259 has **8 intrpt. lines**.
- Sequence of events:
 - 8259 accepts interrupts.
 - 8259 determines priority.
 - 8259 signals 8086 (raises INTR line).
 - CPU Acknowledges.
 - 8259 puts correct vector on data bus.
 - CPU processes interrupt.

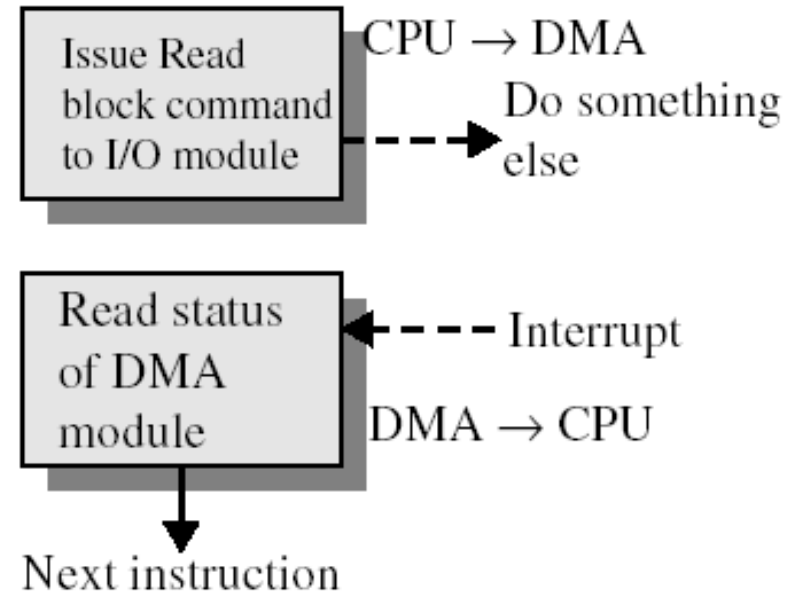


Direct Memory Access (DMA)

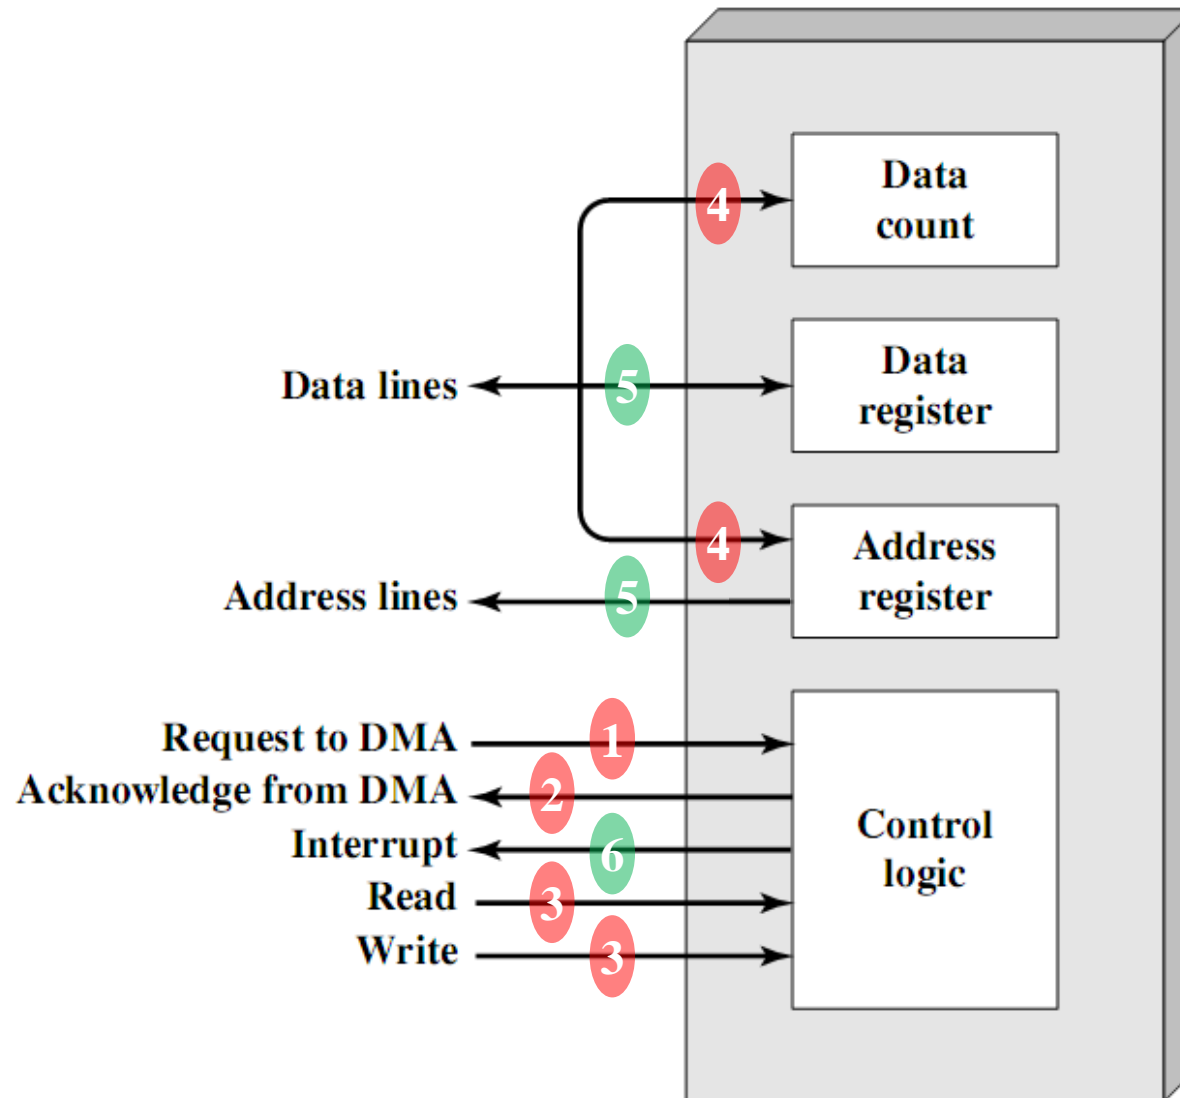
- Interrupt driven and programmed I/O require active CPU intervention.
 - CPU tests and services a device.
 - Transfer rate is limited (depending on CPU availability)!!
 - Many instructions are executed for every I/O.
 - CPU is tied up in managing an I/O transfer!!
- DMA is a more efficient technique (when transferring large volumes of data, i.e., blocks).
 - Additional module on the system bus → DMA controller.
 - DMA controller mimics CPU and takes over the bus to transfer the data with no CPU intervention!

DMA Operation

- CPU **tells** DMA controller:
 - Type of Operation (Rd/Wr).
 - Address of device.
 - Starting address of a data block in memory.
 - Amount of data to be transferred.
- CPU **carries on with other work**.
- DMA controller **performs the transfer**.
- DMA controller **sends interrupt** when finished.



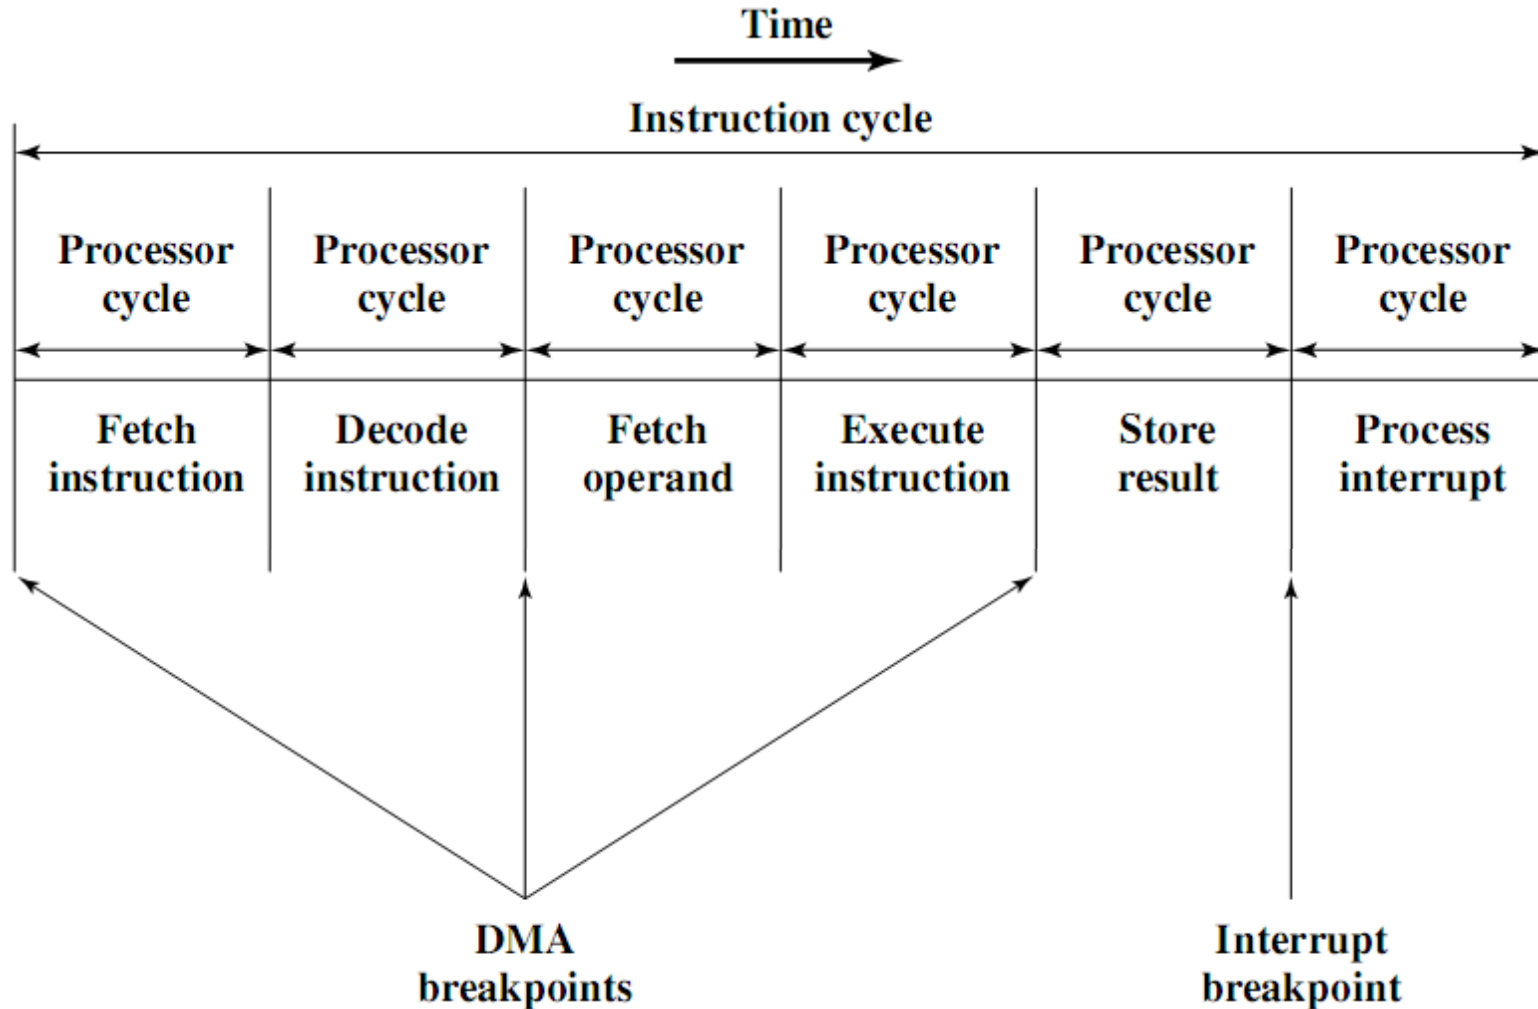
Typical DMA Controller



DMA Transfer modes

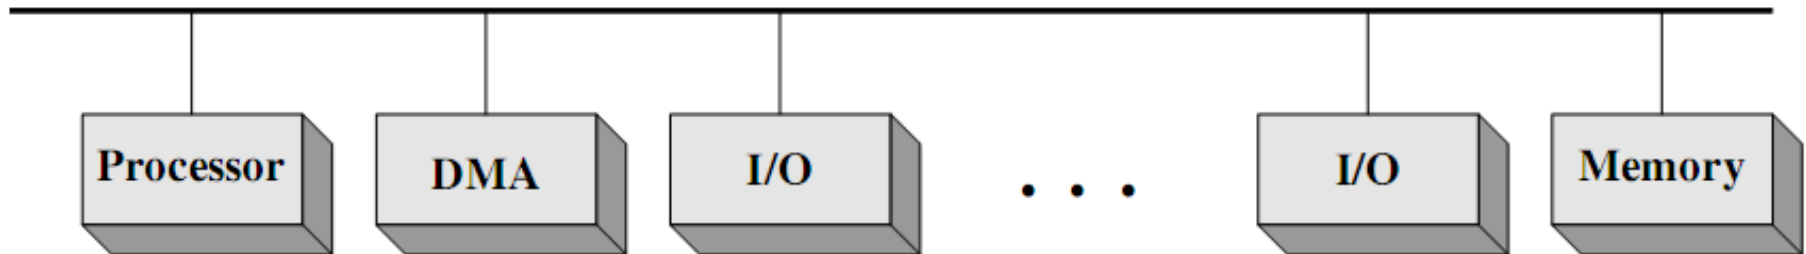
- DMA controller transfers data from/to MM over the system bus.
- DMA controller takes over bus for a bus cycle to transfer data by one of the following techniques:
 - Use bus only when CPU not using it: **transparent** mode.
 - Force CPU to suspend operation temporarily → DMA steals bus cycles from CPU: **cycle stealing** mode.
- Notice this is **not an interrupt!**
 - CPU does not switch context.
- CPU gets suspended just before it accesses bus.
 - i.e. before an operand or data fetch or a data write.
- Slows down CPU. Faster than CPU doing transfer!

DMA and Interrupt Breakpoints During an Instruction Cycle



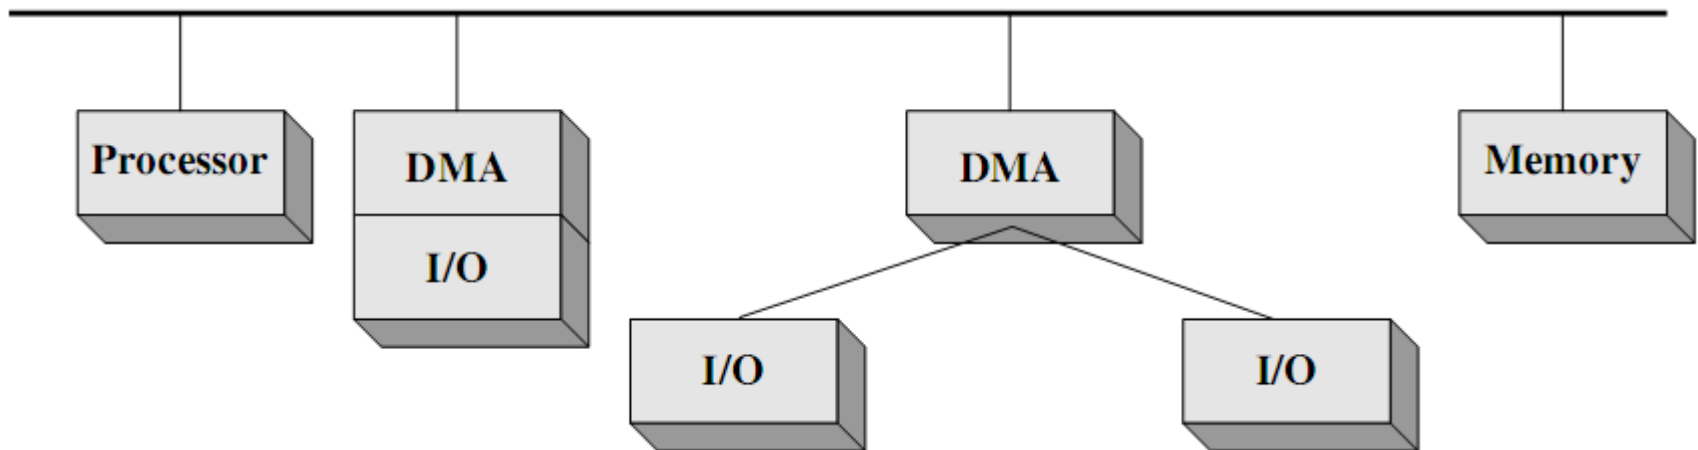
DMA Configurations (1)

- **Features:** single-bus, detached DMA controller.
- DMA module acts as a surrogate processor.
- DMA module uses programmed I/O.
- 👎 • Each transfer uses system **bus twice**.
 - I/O device → DMA controller → memory.
- 👎 • CPU is **suspended twice** per transfer.
- 👍 • DMA controller has **no I/O interfaces**.
- 👍 • **Inexpensive yet inefficient!** 👎



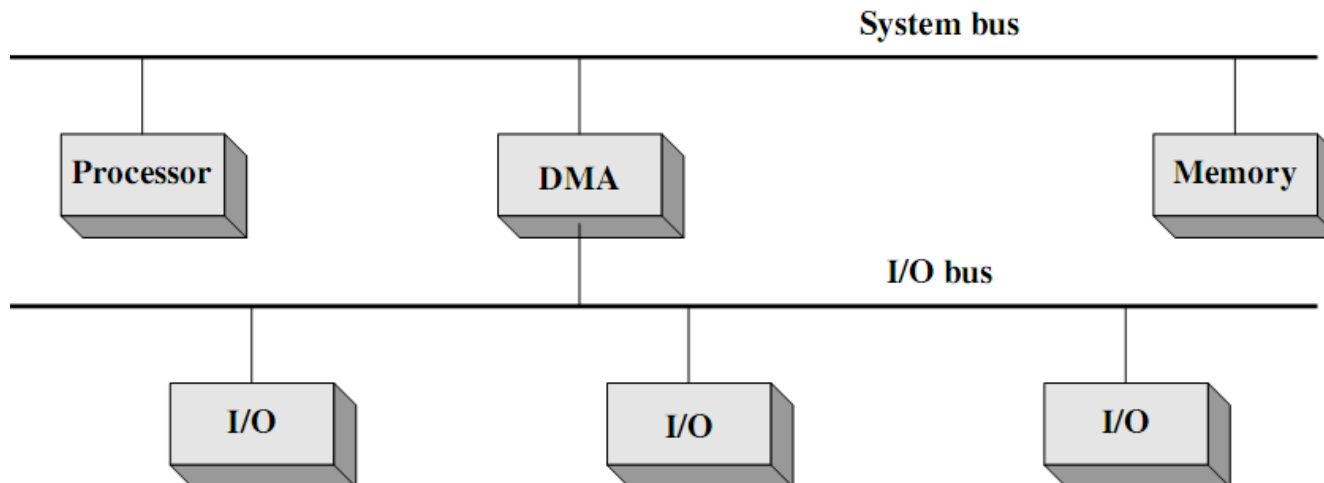
DMA Configurations (2)

- **Features:** single-bus, integrated DMA controller.
- Controller may support more than one device.
- 👍 • Each transfer uses system **bus once**.
 - DMA controller → memory.
- 👍 • CPU is **suspended once** per transfer.
- 👎 • DMA controller has **one or more I/O interfaces**.
- 👍 • **Efficient yet expensive!** 👎



DMA Configurations (3)

- **Features:** separate I/O bus.
 - Connecting all DMA-capable devices.
- 👍 • Each transfer uses system **bus once**.
 - DMA controller → memory
- 👍 • CPU is **suspended once**.
- 👍 • DMA controller has **only one I/O interface**.
- 👍 • **Easily expandable** configuration.



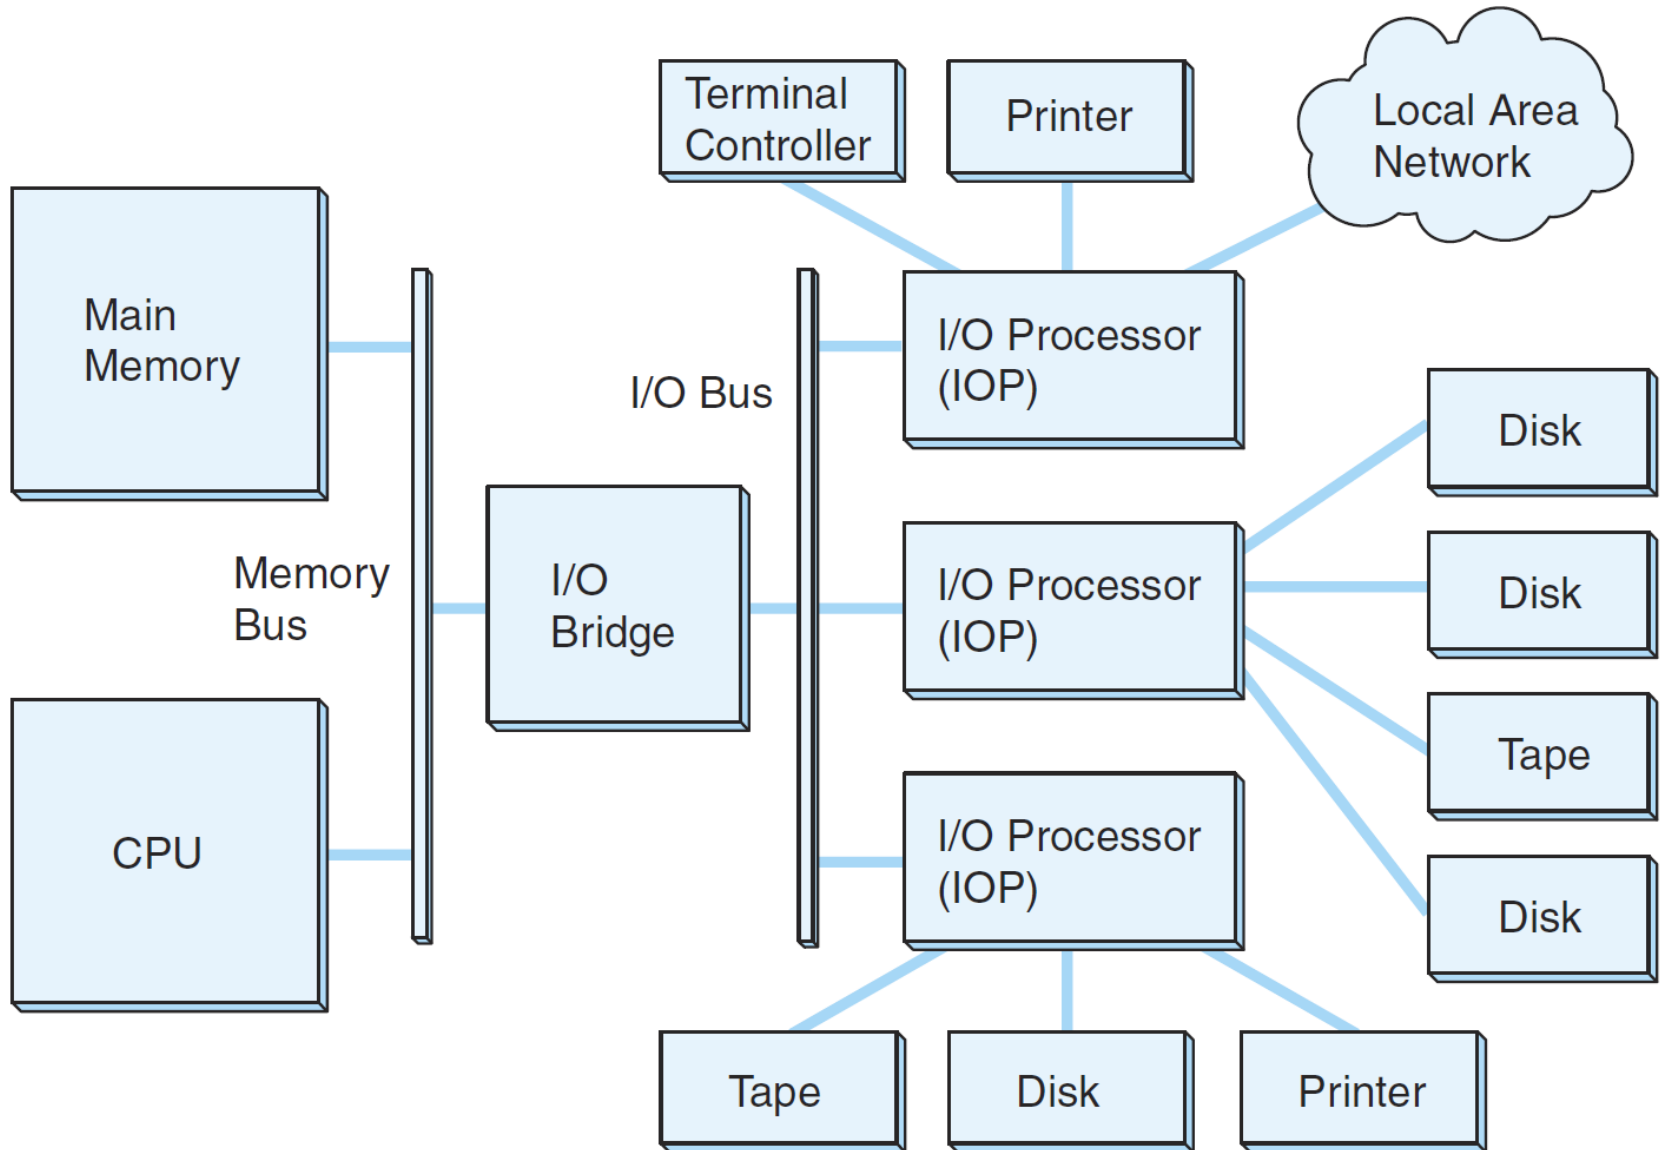
Evolution of I/O function

1. No I/O module.
 - CPU directly controls i/o device.
 2. I/O module **responding** to CPU.
 - Programmed I/O.
 3. I/O module **interrupting** CPU.
 - Interrupt-driven I/O.
 4. I/O module **accessing** memory.
 - Direct-memory access (DMA).
 5. I/O module **executing** program.
 - I/O channel.
 6. I/O module **executing** program from **local memory**.
 - I/O processor, .e.g., GPU.
- **NOTE:** On occasions, no distinction is made between the terms: I/O channel and I/O processor!!!

I/O Channels

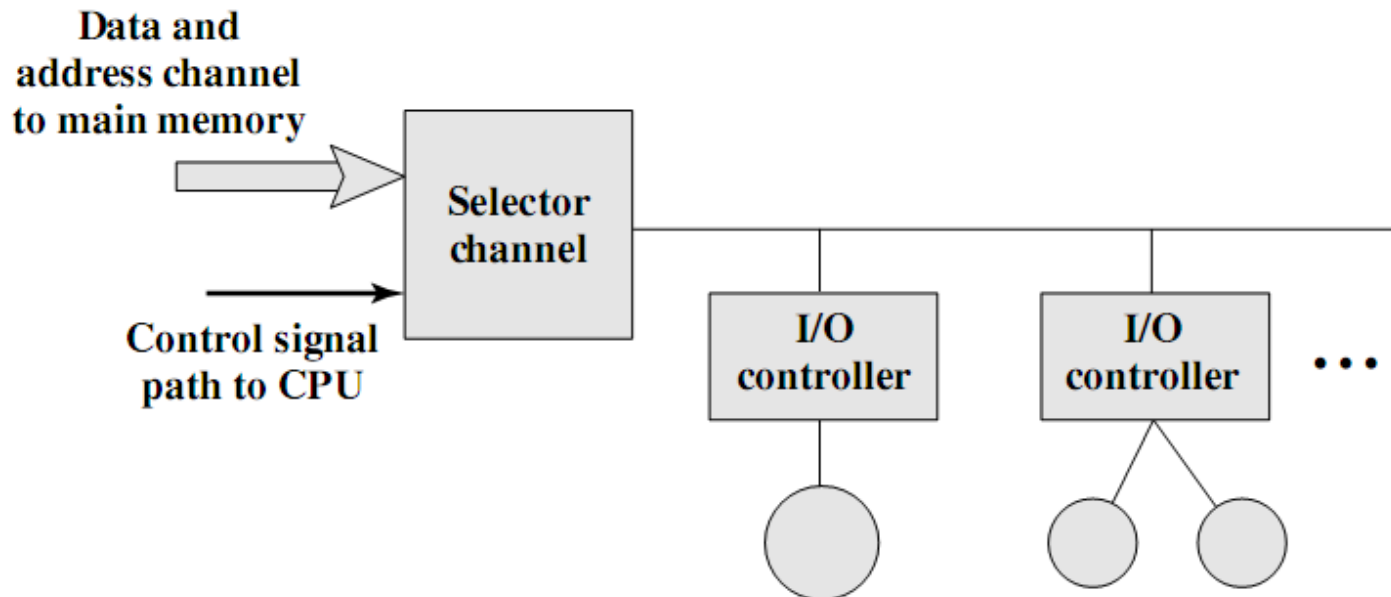
- I/O channels:
 - an extension to the DMA concept
 - able to execute I/O instructions
- I/O channel equipped with **special-purpose processor**, referred to as I/O processor (IOP)!!!!
- CPU **instructs** I/O channel to do the transfer
 - I/O processor fetches and executes I/O program from memory.
- I/O channel does the entire transfer.
- Improves speed
 - Takes load off CPU.
 - Dedicated processor is faster.

A Channel I/O Configuration



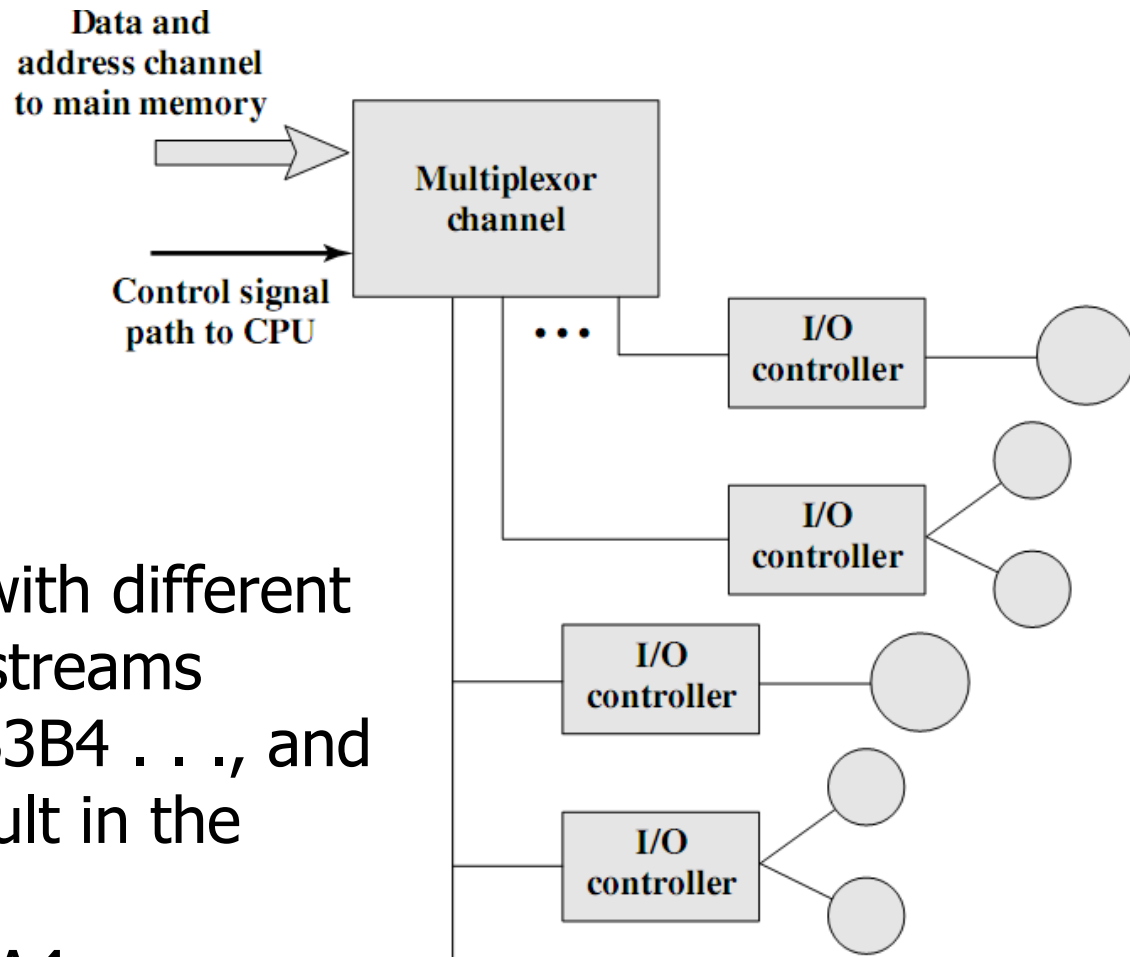
I/O Channel Architecture – Selector

- Channel controls **multiple high-speed** devices.
- At any given time, channel is **dedicated to data transfer with only one** of these devices.
- Each device/set of devices is/are handled by a controller (I/O module).



I/O Channel Architecture – Multiplexor

- Handles **multiple low-speed** devices at the same time.
- A byte multiplexor accepts or transmits characters as fast as possible to multiple devices.
- **Example:** 3 devices with different rates and individual streams A1A2A3A4 ..., B1B2B3B4 . . . , and C1C2C3C4 might result in the character stream A1B1C1A2C2A3B2C3A4 ...



Reading Material

- Stallings, Chapter 7:
 - Pages 232-237
 - Pages 240-243
 - Pages 246-248