


CSE 401


Computer Engineering (2)

هندسة الحاسبات (2)



4th year, Comm. Engineering
Winter 2016

Lecture #10



Dr. Hazem Ibrahim Shehata

Dept. of Computer & Systems Engineering

Credits to Dr. Ahmed Abdul-Monem Ahmed for the slides

Adminstrivia

- Midterm:
 - New date: Thursday, May 5, 2016
 - New time: 12:30pm – 2:00pm
 - Location: classroom #27321 (قاعة 4د)
 - Coverage: lectures #1 → #7
- Assignment #2:
 - For those who emailed me a softcopy, please hand in the original hardcopy this week!!
- Tutorial:
 - To be held on Wednesday at 12:00pm.

Website: <http://hshehata.github.io/courses/zu/cse401/>

Office hours: Monday 11:30am – 12:30pm

Chapter 10. Computer Arithmetic (*Cont.*)

Outline

- Integer Representation
 - Sign-Magnitude, Two's Complement, Biased
- Integer Arithmetic
 - Negation, Addition, Subtraction
 - Multiplication, Division
- Floating-Point Representation
 - IEEE 754
- Floating-Point Arithmetic
 - Addition, Subtraction
 - Multiplication, Division
 - Rounding

Real Numbers

- Numbers with fractions.
- Could be done in pure binary
 - $1001.1010 = 2^3 + 2^0 + 2^{-1} + 2^{-3} = 9.625$
- Where is the binary point?
- Fixed? **0 0 1 0 1 1 0 1 0 0 . 1 1 1 0 1 0**
 - Very large/small numbers cannot be represented.
 - e.g., 0.00000001, 100000000000
 - Fractional part of the quotient in dividing very large numbers will be lost.
- Moving/floating?
 - How do you show where it is?
 - $976,000,000,000,000 = 9.76 \times 10^{14}$
 - $0.000000000000000976 = 9.76 \times 10^{-14}$

**Can do the same with
binary numbers.
What do we need to
store?**

Floating-Point Representation

$$\pm S \times 2^E$$

Sign	Exponent	Significand (Mantissa)
------	----------	------------------------

- The base 2 is the same for all numbers → need not be stored.
- Number is stored in a binary word with 3 fields:
 - Sign: +/–
 - Significand S
 - Exponent E
- **Normal number**: most significant digit of the significand (mantissa) is nonzero → 1 for base 2 (binary).
- What number to store in the significand field? **0.001011**
 - Normal form: 1.011×2^{-3} → Store only 011 in the significand field!
- There is an implicit 1 to the left of the binary point (normalized).
- Exponent indicates place value (floating-point position).

Floating-Point Representation

Biased Exponent


$$\pm S \times 2^E$$



- k-bit unsigned exponent E' ranges from 0 to $2^k - 1$
 - e.g., 8-bit exponent: $0 \leq E' \leq 255$
- The stored exponent E' is a biased exponent
 - $E' = E + (2^{k-1} - 1)$ bias
 - e.g., for 8-bit exponent, $E' = E + 127$
 - $-127 \leq E \leq 128$

• Why?

- Nonnegative floating-point numbers can be treated as unsigned integers for comparison purposes.
- This is not true for 2's comp. or sign-magnitude representations.



Normalization

- FP numbers are usually normalized.
 - i.e., exponent is adjusted so that leading bit (MSB) of mantissa is non-zero, i.e., 1.
 - c.f., Scientific notation where numbers are normalized to give a single digit before the decimal point, e.g. 3.123×10^3 .
- Since the MSB of mantissa is always 1, there is no need to store it!

Floating-Point Examples



<div style="border: 1px solid black; padding: 5px; background-color: #e0ffff; display: inline-block;">1717698.56</div> 1.638125×2^{20} $1.1010001 \times 2^{10100}$	→	<div style="font-size: 3em; vertical-align: middle;">{</div> <div style="display: inline-block; vertical-align: middle;"> Positive → sign bit = 0 $E' = E + 127 = 10100 + 1111111 = 10010011$ Mantissa = 1010001 0000000000000000 </div>	<div style="border: 1px solid black; padding: 5px; background-color: yellow; text-align: center;"> 0 10010011 101000100000000000000000 </div>
--	---	---	---


<div style="border: 1px solid black; padding: 5px; background-color: #e0ffff; display: inline-block;">-1717698.56</div> -1.638125×2^{20} $-1.1010001 \times 2^{10100}$	→	<div style="font-size: 3em; vertical-align: middle;">{</div> <div style="display: inline-block; vertical-align: middle;"> Negative → sign bit = 1 $E' = E + 127 = 10100 + 1111111 = 10010011$ Mantissa = 1010001 0000000000000000 </div>	<div style="border: 1px solid black; padding: 5px; background-color: yellow; text-align: center;"> 1 10010011 101000100000000000000000 </div>
---	---	---	---

<div style="border: 1px solid black; padding: 5px; background-color: #e0ffff; display: inline-block;">...</div> 1.638125×2^{-20} $1.1010001 \times 2^{-10100}$	→	<div style="font-size: 3em; vertical-align: middle;">{</div> <div style="display: inline-block; vertical-align: middle;"> Positive → sign bit = 0 $E' = E + 127 = -10100 + 1111111 = 01101011$ Mantissa = 1010001 0000000000000000 </div>	<div style="border: 1px solid black; padding: 5px; background-color: yellow; text-align: center;"> 0 01101011 101000100000000000000000 </div>
---	---	--	---


FP Ranges (32-bit)

- 32-bit FP number, 8-bit exponent, 23-bit mantissa.
- Largest +ve number $(2 - 2^{-23}) \times 2^{128}$
 - Largest true exponent: 128 $0.111...11$
 - Largest mantissa: $1 + (1 - 2^{-23}) = 2 - 2^{-23}$
- Smallest +ve number 2^{-127}
 - Smallest true exponent: -127
 - Smallest mantissa: 1
- Smallest -ve number $-(2 - 2^{-23}) \times 2^{128}$
- Largest -ve number -2^{-127}
- Accuracy
 - The effect of changing LSB of mantissa.
 - 23-bit mantissa $2^{-23} \approx 1.2 \times 10^{-7}$
 - About 6 decimal places.

Expressible Numbers (32-bit)



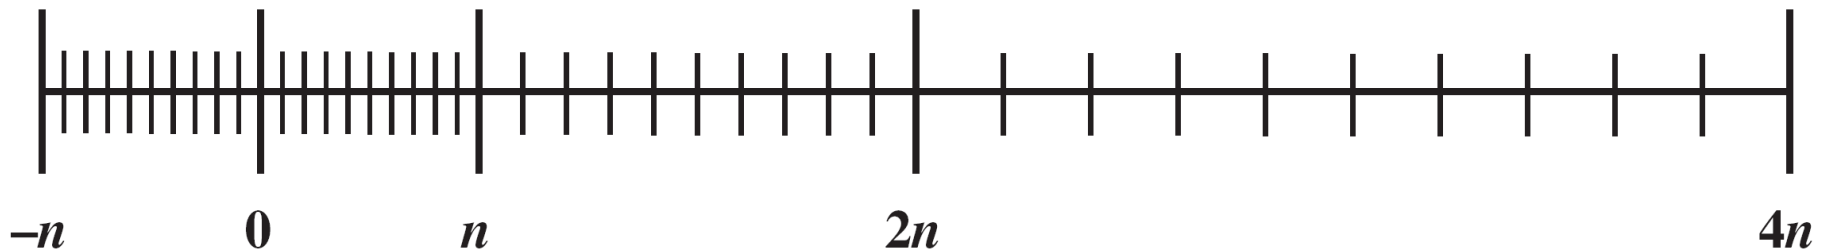
(a) Two's Complement Integers



(b) Floating-Point Numbers

Density of Floating Point Numbers

- 32-bit FP number $\rightarrow 2^{32}$ different values represented.
- No more individual values are represented with floating-point numbers. Numbers are just spread out.
- Numbers represented in the FP representation are not spaced evenly along the line number. Why?
- Range-precision tradeoff
 - More bits for exponent \rightarrow wider range & less precision
 - **Reason**: there is a fixed number of values that can be represented!
 - To increase both range and precision \rightarrow use more bits!!!

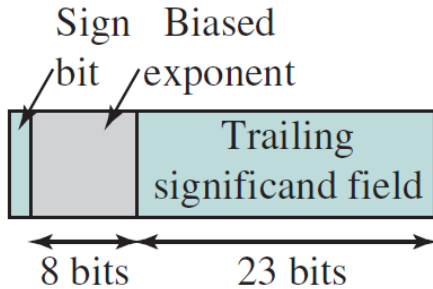


IEEE 754

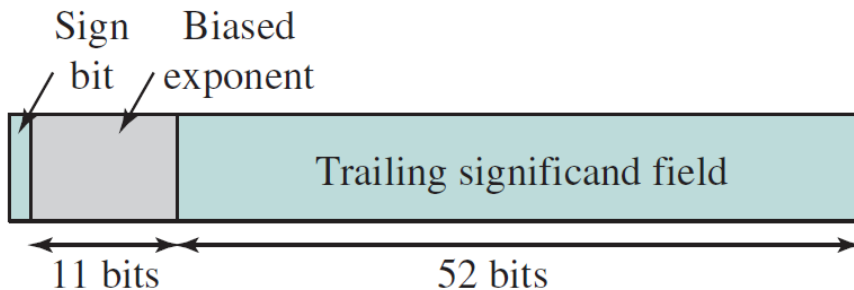
- Standard for floating-point representation.
- Adopted 1985 and revised 2008.
- IEEE 754-2008 defines many FP formats for different purposes:

Format	Format Type		
	Arithmetic Format	Basic Format	Interchange Format
binary16			X
binary32	X	X	X
binary64	X	X	X
binary128	X	X	X
binary{k} ($k = n \times 32$ for $n > 4$)	X		X
decimal64	X	X	X
decimal128	X	X	X
decimal{k} ($k = n \times 32$ for $n > 4$)	X		X
extended precision	X		
extendable precision	X		

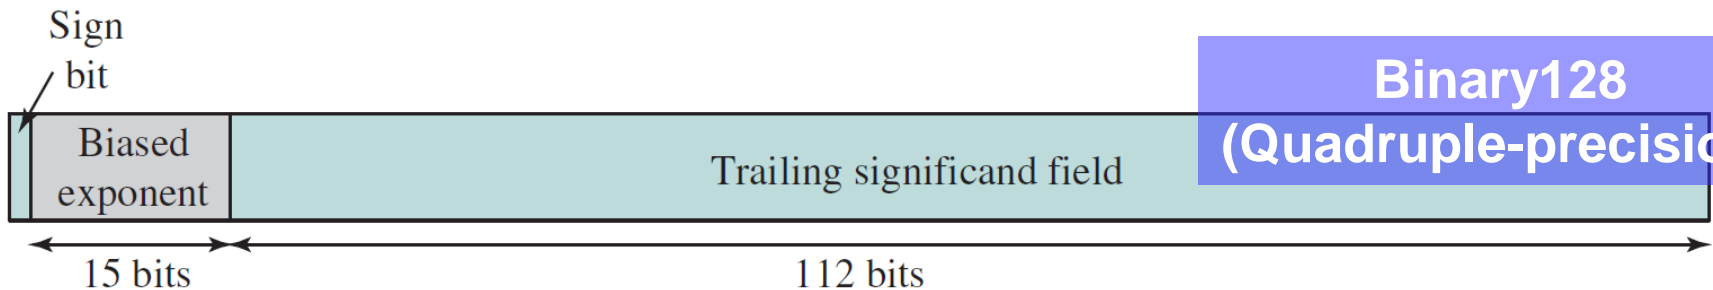
IEEE 754 - Binary³²/⁶⁴/¹²⁸ Formats



Binary32
(Single-precision)



Binary64
(Double-precision)



Binary128
(Quadruple-precision)

IEEE 754 - Binary32/64/128 Interpretations

	Sign	Biased Exponent	Fraction	Value
positive zero	0	0	0	0
negative zero	1	0	0	-0
plus infinity	0	all 1s	0	∞
minus infinity	1	all 1s	0	$-\infty$
quiet NaN	0 or 1	all 1s	$\neq 0$; first bit = 1	qNaN
signaling NaN	0 or 1	all 1s	$\neq 0$; first bit = 0	sNaN
positive normal nonzero	0	$0 < e < 255$	f	$2^{e-127}(1.f)$
negative normal nonzero	1	$0 < e < 255$	f	$-2^{e-127}(1.f)$
positive subnormal	0	0	$f \neq 0$	$2^{-126}(0.f)$
negative subnormal	1	0	$f \neq 0$	$-2^{-126}(0.f)$
positive normal nonzero	0	$0 < e < 2047$	f	$2^{e-1023}(1.f)$
negative normal nonzero	1	$0 < e < 2047$	f	$-2^{e-1023}(1.f)$
positive subnormal	0	0	$f \neq 0$	$2^{-1022}(0.f)$
negative subnormal	1	0	$f \neq 0$	$-2^{-1022}(0.f)$
positive normal nonzero	0	$0 < e < 32767$	f	$2^{e-16383}(1.f)$
negative normal nonzero	1	$0 < e < 32767$	f	$-2^{e-16383}(1.f)$
positive subnormal	0	0	$f \neq 0$	$2^{-16382}(0.f)$
negative subnormal	1	0	$f \neq 0$	$-2^{-16382}(0.f)$

32

64

128

IEEE 754 - Binary³²/⁶⁴/¹²⁸ Parameters

Parameter	Format		
	Binary32	Binary64	Binary128
Storage width (bits)	32	64	128
Exponent width (bits)	8	11	15
Exponent bias	127	1023	16383
Maximum exponent	127	1023	16383
Minimum exponent	-126	-1022	-16382
Approx normal number range (base 10)	$10^{-38}, 10^{+38}$	$10^{-308}, 10^{+308}$	$10^{-4932}, 10^{+4932}$
Trailing significand width (bits)*	23	52	112
Number of exponents	254	2046	32766
Number of fractions	2^{23}	2^{52}	2^{112}
Number of values	1.98×2^{31}	1.99×2^{63}	1.99×2^{128}
Smallest positive normal number	2^{-126}	2^{-1022}	2^{-16362}
Largest positive normal number	$2^{128} - 2^{104}$	$2^{1024} - 2^{971}$	$2^{16384} - 2^{16271}$
Smallest subnormal magnitude	2^{-149}	2^{-1074}	2^{-16494}

Note: *not including implied bit and not including sign bit

IEEE 754 - NaNs

- NaN:

- Symbolic entity encoded in FP format
- Types: Signaling (sNaN) or Quiet (qNaN)
- Both types have the same format:

S = 0 or 1

E = 1111...11

F ≠ 0000..00

- F distinguishes between the two types:

— F=0xxxx..xx → sNaN, F=1xxxx..xx → qNaN

- Signaling NaN:

- Signals an invalid operation exception whenever it appears as an operand. Ex.: uninitialized variables


- Quite NaN:

- Propagates without signaling exceptions.


IEEE 754 - Quiet NaN

Operation	Quiet NaN Produced By
Any	Any operation on a signaling NaN
Add or subtract	Magnitude subtraction of infinities: $(+\infty) + (-\infty)$ $(-\infty) + (+\infty)$ $(+\infty) - (+\infty)$ $(-\infty) - (-\infty)$
Multiply	$0 \times \infty$
Division	$\frac{0}{0}$ or $\frac{\infty}{\infty}$
Remainder	$x \text{ REM } 0$ or $\infty \text{ REM } y$
Square root	\sqrt{x} , where $x < 0$

IEEE 754 - Effect of Subnormal Numbers



(a) 32-Bit format without subnormal numbers

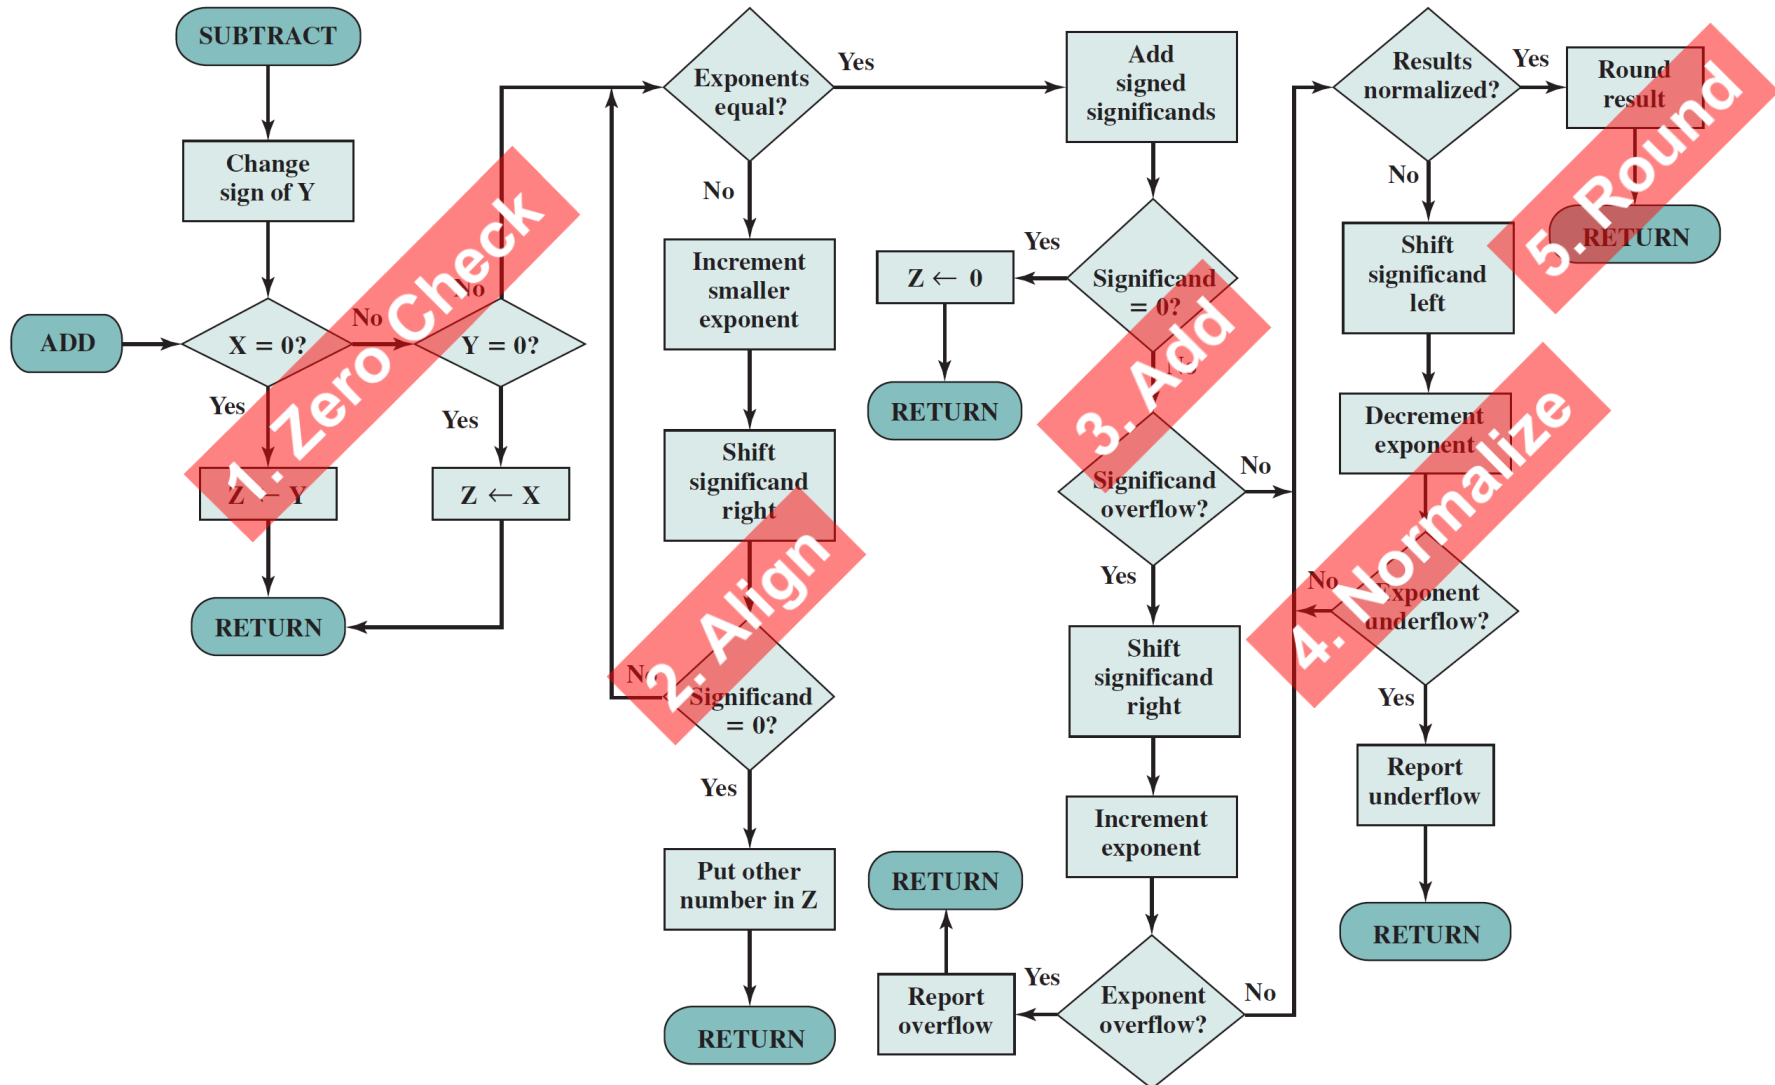


(b) 32-Bit format with subnormal numbers

FP Arithmetic +/-

- Algorithm:
 1. Check for zeros.
 2. Align significands (adjusting exponents).
 3. Add or subtract significands.
 4. Normalize result.
 5. Round result.

FP Addition & Subtraction Flowchart



Reading Material

- Stallings, Chapter 10:
 - Pages 341-352
 - Pages 356-358