

CSE 411: Artificial Intelligence (Elective Course #6)

**400 Level, Mechatronics Engineering
2nd Term 2016/2017, Lecture #5**

Hazem Shehata

Dept. of Computer & Systems Engineering
Zagazig University

March 27th, 2017

Credits to Dr. Mohamed El Abd for the slides

Adminstrivia

Notes

- Assignment #2:
 - To be released on Thursday (Due 1 week after).
 - Constructing search trees on paper.

Course Info:

- Website: <http://hshehata.github.io/courses/zu/cse411/>
- Office hours: Sunday 11:30am - 12:30pm

Adminstrivia

Notes

- Assignment #2:
 - To be released on Thursday (Due 1 week after).
 - Constructing search trees on paper.
- Assignment #3:
 - To be released next week.
 - Implementing search algorithms in Python.

Course Info:

- Website: <http://hshehata.github.io/courses/zu/cse411/>
- Office hours: Sunday 11:30am - 12:30pm

Search

Types of search algorithms

- **Uninformed Search:**

Only has the information provided by the problem formulation (initial state, available actions, transition model, goal test, and step/path cost).

Types of search algorithms

- **Uninformed Search:**

Only has the information provided by the problem formulation (initial state, available actions, transition model, goal test, and step/path cost).

- Algorithms: **BFS, DFS, DLS, IDS, UCS.**

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Types of search algorithms

- **Uninformed Search:**

Only has the information provided by the problem formulation (initial state, available actions, transition model, goal test, and step/path cost).

- Algorithms: **BFS, DFS, DLS, IDS, UCS.**

- **Informed Search:**

Has additional information that allows it to judge the promise of an action, *i.e.*, the estimated cost from a state to a goal.

Types of search algorithms

- **Uninformed Search:**

Only has the information provided by the problem formulation (initial state, available actions, transition model, goal test, and step/path cost).

- Algorithms: **BFS, DFS, DLS, IDS, UCS.**

- **Informed Search:**

Has additional information that allows it to judge the promise of an action, *i.e.*, the estimated cost from a state to a goal.

- Algorithms: **GBFS, A*.**

Outline

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

1 Informed Search

2 Requirements & Reading Material

Outline

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

1 Informed Search

2 Requirements & Reading Material

Informed Search

Introduction

- Uninformed search algorithms are systematic but inefficient.

Informed Search

Introduction

- Uninformed search algorithms are systematic but inefficient.
- They just repeat a cycle of:
 - Choosing the next node to expand.
 - Check whether it is the goal.
 - If it is not, expand the node.

Informed Search

Introduction

- Uninformed search algorithms are systematic but inefficient.
- They just repeat a cycle of:
 - Choosing the next node to expand.
 - Check whether it is the goal.
 - If it is not, expand the node.
- Nodes are chosen in a specific order (imposed by a search strategy), which is not necessarily the "best" order in terms of getting closer/faster to the goal.

Informed Search

Introduction

- Would like to use additional knowledge so that "better" nodes are expanded/explored first.

Informed Search

Introduction

- Would like to use additional knowledge so that "better" nodes are expanded/explored first.
- In terms of pseudo-code, we want to order the frontier so that "better" nodes get popped first.

Informed Search

Introduction

- Would like to use additional knowledge so that "better" nodes are expanded/explored first.
- In terms of pseudo-code, we want to order the frontier so that "better" nodes get popped first.
- We will introduce an **evaluation function** $f(n)$ that indicates the desirability of considering node n next for exploration and expansion.

Informed Search

Introduction

- Would like to use additional knowledge so that "better" nodes are expanded/explored first.
- In terms of pseudo-code, we want to order the frontier so that "better" nodes get popped first.
- We will introduce an **evaluation function** $f(n)$ that indicates the desirability of considering node n next for exploration and expansion.
- Nodes with a better $f(n)$ are always considered first. This approach is called **best-first search**.

Informed Search

Introduction

- Would like to use additional knowledge so that "better" nodes are expanded/explored first.
- In terms of pseudo-code, we want to order the frontier so that "better" nodes get popped first.
- We will introduce an **evaluation function** $f(n)$ that indicates the desirability of considering node n next for exploration and expansion.
- Nodes with a better $f(n)$ are always considered first. This approach is called **best-first search**.
- How should we compute $f(n)$?

Informed Search

Introduction

- Uniform-cost search orders the frontier according to the path cost $g(n)$:
 - Path cost is distance from root to state n .

Informed Search

Introduction

- Uniform-cost search orders the frontier according to the path cost $g(n)$:
 - Path cost is distance from root to state n .
- So, uniform-cost search uses an evaluation function $f(n) = g(n)$.

Informed Search

Introduction

- Uniform-cost search orders the frontier according to the path cost $g(n)$:
 - Path cost is distance from root to state n .
- So, uniform-cost search uses an evaluation function $f(n) = g(n)$.
- The path cost $g(n)$ only accounts for cost to reach n . Hence, uniform-cost search is not goal directed/oriented.

Informed Search

Introduction

- Would also like to consider cost from node n to the goal.

Informed Search

Introduction

- Would also like to consider cost from node n to the goal.
- More generally:
We want the search algorithm to be able to *estimate* the path cost from the current node to the goal.

Informed Search

Introduction

- Would also like to consider cost from node n to the goal.
- More generally:
We want the search algorithm to be able to *estimate* the path cost from the current node to the goal.
- This estimate is called a **heuristic function**.

Informed Search

Introduction

- Would also like to consider cost from node n to the goal.
- More generally:
We want the search algorithm to be able to *estimate* the path cost from the current node to the goal.
- This estimate is called a **heuristic function**.
- Cannot be done based on problem formulation:
Need to add additional information.

Informed Search

Introduction

- Path cost $g(n)$ is an exact value.

Informed Search

Introduction

- Path cost $g(n)$ is an exact value.
- Heuristic function $h(n)$:
 - $h(n)$: estimated cost from node n to goal.
 - $h(n1) < h(n2)$ means it is probably cheaper to get to the goal from $n1$.
 - $h(goal) = 0$.

Informed Search

Introduction

- Path cost $g(n)$ is an exact value.
- Heuristic function $h(n)$:
 - $h(n)$: estimated cost from node n to goal.
 - $h(n1) < h(n2)$ means it is probably cheaper to get to the goal from $n1$.
 - $h(goal) = 0$.
- Evaluation function $f(n)$:
 - $f(n) = g(n)$: Uniform-cost search.
 - $f(n) = h(n)$: Greedy best-first search.
 - $f(n) = g(n) + h(n)$: A* search.

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Informed Search

Greedy best-first search

- GBFS algorithm always explores and expands the node judged to be closest to goal.

Informed Search

Greedy best-first search

- GBFS algorithm always explores and expands the node judged to be closest to goal.
- It uses $f(n) = h(n)$ and ignores the path cost $g(n)$ entirely.

Informed Search

Greedy best-first search

- GBFS algorithm always explores and expands the node judged to be closest to goal.
- It uses $f(n) = h(n)$ and ignores the path cost $g(n)$ entirely.
- Consider it to be the complement of uniform-cost search.

Informed Search

Greedy best-first search

- GBFS algorithm always explores and expands the node judged to be closest to goal.
- It uses $f(n) = h(n)$ and ignores the path cost $g(n)$ entirely.
- Consider it to be the complement of uniform-cost search.
- GBFS algorithm is identical to UNIFORM-COST-SEARCH except that h is used instead of g .

Search Algorithms

Greedy best-first search algorithm (tree version)

```
function GREEDY-BEST-FIRST-SEARCH(problem) returns a solution, or failure
  node ← node with STATE=problem.INITIAL-STATE, PATH-COST=0
  frontier ← priority queue ordered by heuristic h, with node as only element
  loop do
    if frontier.EMPTY?() then return failure
    node ← frontier.POP()          /* choose lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return node.SOLUTION()
    for each action in problem.ACTIONS(node.STATE) do
      child ← node.CHILD-NODE(problem, action)
      frontier.INSERT(child, h(child))
```

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Search Algorithms

Greedy best-first search algorithm (graph version)

```
function GREEDY-BEST-FIRST-SEARCH(problem) returns a solution, or failure
  node ← node with STATE=problem.INITIAL-STATE, PATH-COST=0
  frontier ← priority queue ordered by heuristic h, with node as only element
  explored ← an empty set
  loop do
    if frontier.EMPTY?() then return failure
    node ← frontier.POP()           /* choose lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return node.SOLUTION()
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child ← node.CHILD-NODE(problem, action)
      if child.STATE is not in explored and not in frontier then
        frontier.INSERT(child, h(child))
      else if child.STATE is in frontier with higher h-value then
        replace that frontier node with child
```

Outline

Informed
Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements
& Reading
Material

Informed Search

GBFS example - Romania map

Find a route from Arad to Bucharest.

Outline

Informed Search


Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material



straight-line distances to Bucharest

| | |
|----------------|-----|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Drobeta | 242 |
| Eforie | 161 |
| Fagaras | 176 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 100 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |

Informed Search

GBFS example - Romania map

Find a route from Arad to Bucharest.


(a) The initial state




(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Fagaras



Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Informed Search

Greedy best-first search

- Not optimal.

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Informed Search

Greedy best-first search

- Not complete (unless m is finite; uncommon in trees).
 - May get stuck in an infinite loop (e.g., reaching a deadend through a reversible action).
- Not optimal.

Informed Search

Greedy best-first search

- Not complete (unless m is finite; uncommon in trees).
 - May get stuck in an infinite loop (e.g., reaching a deadend through a reversible action).
- Not optimal.
- Time complexity = $O(b^m)$.

Informed Search

Greedy best-first search


- Not complete (unless m is finite; uncommon in trees).
 - May get stuck in an infinite loop (e.g., reaching a deadend through a reversible action).
- Not optimal.
- Time complexity = $O(b^m)$.
- Space complexity = $O(b^m)$.

Informed Search

Greedy best-first search

GBFS properties:

- Upper-bound case: goal is last node of the tree:
 - Number of nodes generated: b nodes for each node of m levels (entire tree).
 - Time and space complexity: all generated nodes $O(b^m)$.



Outline

Informed
Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements
& Reading
Material

Informed Search

A* search

- Uniform cost search orders the queue according to the path cost $g(n)$:
 - Optimal, complete, but inefficient in time and space.

Informed Search

A* search

- Uniform cost search orders the queue according to the path cost $g(n)$:
 - Optimal, complete, but inefficient in time and space.
- Greedy best first search orders the queue using the heuristic cost $h(n)$:
 - Not optimal, not complete but efficient and directed (with good heuristic).

Informed Search

A* search

- Idea behind A* is to combine the two strategies:

Informed Search

A* search

- Idea behind A* is to combine the two strategies:
 - Use an evaluation function $f(n) = g(n) + h(n)$ to order the nodes to be explored.

Informed Search

A* search

- Idea behind A* is to combine the two strategies:
 - Use an evaluation function $f(n) = g(n) + h(n)$ to order the nodes to be explored.
 - $f(n)$ measures the cheapest total estimated cost from the initial state to the goal state passing through the current state n .

Informed Search

A* search

- Idea behind A* is to combine the two strategies:
 - Use an evaluation function $f(n) = g(n) + h(n)$ to order the nodes to be explored.
 - $f(n)$ measures the cheapest total estimated cost from the initial state to the goal state passing through the current state n .
- The resulting search is both optimal and complete assuming certain conditions on the heuristic cost $h(n)$.

Informed Search

A* search

- Idea behind A* is to combine the two strategies:
 - Use an evaluation function $f(n) = g(n) + h(n)$ to order the nodes to be explored.
 - $f(n)$ measures the cheapest total estimated cost from the initial state to the goal state passing through the current state n .
- The resulting search is both optimal and complete assuming certain conditions on the heuristic cost $h(n)$.
- A* algorithm is identical to UNIFORM-COST-SEARCH except that $g + h$ is used instead of g .

Search Algorithms

A* search algorithm (tree version)

```
function A-STAR-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  node with STATE=problem.INITIAL-STATE, PATH-COST=0
  frontier  $\leftarrow$  priority queue ordered by evaluation f, with node as only element
  loop do
    if frontier.EMPTY?() then return failure
    node  $\leftarrow$  frontier.POP() /* choose lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return node.SOLUTION()
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  node.CHILD-NODE(problem, action)
      frontier.INSERT(child, f(child))
```

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Search Algorithms

A* search algorithm (graph version)

```
function A-STAR-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  node with STATE=problem.INITIAL-STATE, PATH-COST=0
  frontier  $\leftarrow$  priority queue ordered by evaluation  $f$ , with node as only element
  explored  $\leftarrow$  an empty set
  loop do
    if frontier.EMPTY?() then return failure
    node  $\leftarrow$  frontier.POP()          /* choose lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return node.SOLUTION()
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  node.CHILD-NODE(problem, action)
      if child.STATE is not in explored and not in frontier then
        frontier.INSERT(child, evaluation  $f$ (child))
      else if child.STATE is in frontier with higher evaluation  $f$ -value then
        replace that frontier node with child
```

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Informed Search

A* example - Romania map

Find a route from Arad to Bucharest.

Outline

Informed Search

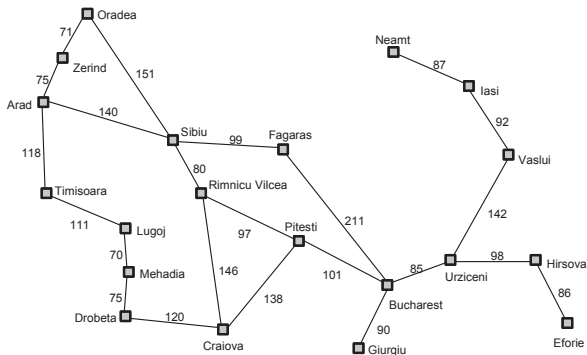
Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material



straight-line distances to Bucharest

| | |
|----------------|-----|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Drobeta | 242 |
| Eforie | 161 |
| Fagaras | 176 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 100 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |

Informed Search

A* example - Romania map

Find a route from Arad to Bucharest.


(a) The initial state




(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Rimnicu Vilcea



Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

Informed Search

A* example - Romania map

Find a route from Arad to Bucharest.

Outline

Informed Search

Introduction


Greedy best-first

A*


Heuristic Functions

Requirements & Reading Material

(e) After expanding Fagaras



(f) After expanding Pitesti



Informed Search

A* search

- Optimal, given an **admissible heuristic**.

Informed Search

A* search

- Complete (if b and ε are finite).
 - Reason: number of nodes with cost $\leq C^*$ is finite.
- Optimal, given an **admissible heuristic**.

Informed Search

A* search

- Complete (if b and ε are finite).
 - Reason: number of nodes with cost $\leq C^*$ is finite.
- Optimal, given an **admissible heuristic**.
- Time and space complexity: not straightforward!
 - Number of nodes explored depends on the difference between h and h^* (true cost).
 - If $h = h^*$, A* expands only the nodes on the optimal solution path(s).
 - If $h = 0$, A* consumes as much (time/space) resources as UCS.

Informed Search

A* search

- A heuristic $h(n)$ is admissible if for every node n :

$$h(n) \leq h^*(n),$$

where $h^*(n)$ is the true cost to reach goal state from n .

Informed Search

A* search

- A heuristic $h(n)$ is admissible if for every node n :
$$h(n) \leq h^*(n),$$

where $h^*(n)$ is the true cost to reach goal state from n .
- An admissible heuristic function $h(n)$ **never overestimates the true cost from n to goal**.

Informed Search

A* search

- A heuristic $h(n)$ is admissible if for every node n :
$$h(n) \leq h^*(n),$$

where $h^*(n)$ is the true cost to reach goal state from n .
- An admissible heuristic function $h(n)$ **never overestimates the true cost from n to goal**.
- If $h(n)$ never overestimates, then $f(n)$ never overestimates the true cost to the goal through node n .

Informed Search

A* search

- Proving that "admissibility guarantees optimality":

Outline

Informed Search


Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material



Informed Search

A* search

- Proving that "admissibility guarantees optimality":
 - Suppose an A* search (that uses an admissible heuristic h) finds a goal G whose path cost of is C .
 - $h(G) = 0$, and hence: $f(G) = g(G) = C$
 - Let G' be another goal node whose path cost is C' .
 - Frontier contains a node n on the optimal path to G' .
 - h is admissible, and hence: $f(n) = g(n) + h(n) \leq C'$
 - Node G has been popped from frontier before node n .
 - $f(G) \leq f(n)$, and hence: $C \leq C'$.
 - Using an admissible heuristic, A* would always discover the lowest-cost (*i.e.*, optimal) solution.

Informed Search

A* search

- In graph-search algorithms, a new node is discarded if it's already in the *explored* set.

Informed Search

A* search

- In graph-search algorithms, a new node is discarded if it's already in the *explored* set.
- If the new node has a better path cost (g) than the old node, a shorter path has been ignored.

Informed Search

A* search

- In graph-search algorithms, a new node is discarded if it's already in the *explored* set.
- If the new node has a better path cost (g) than the old node, a shorter path has been ignored.
- This means that the graph version of A* search is not optimal any more!

Informed Search

A* search


- In graph-search algorithms, a new node is discarded if it's already in the *explored* set.
- If the new node has a better path cost (g) than the old node, a shorter path has been ignored.
- This means that the graph version of A* search is not optimal any more!
- To ensure optimality, A* must use a **consistent** heuristic function h .
 - Consistency is a stronger condition than admissibility.

Informed Search

A* search

- **Consistency** (monotonicity) means that for every node n and child node n' reachable from n by action a , the estimated cost $h(n)$ is never greater than the estimated cost $h(n')$ plus the step cost of getting to n' :

$$h(n) \leq h(n') + c(n, a, n')$$



Informed Search

A* search

- A consistent heuristic function $h(n)$ **never overestimates the true cost of action a taken from n to n' :**

$$h(n) - h(n') \leq c(n, a, n')$$

Informed Search

A* search

- A consistent heuristic function $h(n)$ **never overestimates the true cost of action a taken from n to n' :**

$$h(n) - h(n') \leq c(n, a, n')$$

- With some mathematical arrangements:

$$\begin{aligned} f(n) &= h(n) + g(n) \\ &\leq h(n') + c(n, a, n') + g(n) \\ &\leq h(n') + g(n') \\ &\leq f(n') \end{aligned}$$

So, $f(n)$ never decreases as we approach the goal.

Informed Search

A* search

- A consistent heuristic function $h(n)$ **never overestimates the true cost of action a taken from n to n'** :

$$h(n) - h(n') \leq c(n, a, n')$$

- With some mathematical arrangements:

$$\begin{aligned} f(n) &= h(n) + g(n) \\ &\leq h(n') + c(n, a, n') + g(n) \\ &\leq h(n') + g(n') \\ &\leq f(n') \end{aligned}$$

So, $f(n)$ never decreases as we approach the goal.

- Consistency guarantees that states are always visited by the cheapest path first; no need to check if subsequent paths are better than the first.

Informed Search

A* example - consistency

Check whether h is consistent, and perform A* search.

Outline

Informed Search


Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material



Informed Search

Introduction

- Informed search algorithms use additional knowledge about the problem to direct search toward goal(s).

Informed Search

Introduction

- Informed search algorithms use additional knowledge about the problem to direct search toward goal(s).
- Such additional knowledge takes the form of a **heuristic function** h .
 - $h(n)$: *estimated* path cost from node n to closest goal.

Informed Search

Introduction

- Informed search algorithms use additional knowledge about the problem to direct search toward goal(s).
- Such additional knowledge takes the form of a **heuristic function** h .
 - $h(n)$: *estimated* path cost from node n to closest goal.
- Frontier nodes are ordered using an **evaluation function** f defined in terms of h .
 - GBFS: $f(n) = h(n)$.
 - A*: $f(n) = h(n) + g(n)$.

Informed Search


Introduction

- Informed search algorithms use additional knowledge about the problem to direct search toward goal(s).
- Such additional knowledge takes the form of a **heuristic function** h .
 - $h(n)$: *estimated* path cost from node n to closest goal.
- Frontier nodes are ordered using an **evaluation function** f defined in terms of h .
 - GBFS: $f(n) = h(n)$.
 - A*: $f(n) = h(n) + g(n)$.
- A* has the advantage of being optimal given that:
 - h is admissible (tree version).
 - h is consistent (graph version).

Informed Search

Effect of heuristic functions


- Uniform-cost search expands in circular cost contours ($h(n) = 0$).



Informed Search

Effect of heuristic functions

- Uniform-cost search expands in circular cost contours ($h(n) = 0$).
- A* search elongates & rotates contours towards goal:
 - More narrow/elongated, the better $h(n)$ is. More directed.



Informed Search

Quality of heuristic function

- One heuristic function might be better than another for a given problem.

Informed Search

Quality of heuristic function

- One heuristic function might be better than another for a given problem.

- Informedness:**

For two admissible heuristic functions, h_1 and h_2 :

if $h_2(n) \geq h_1(n)$

$h_2(n)$ is more informed than $h_1(n)$

Alternatively we say that $h_2(n)$ dominates $h_1(n)$.

Informed Search

Quality of heuristic function

- One heuristic function might be better than another for a given problem.

- **Informedness:**

For two admissible heuristic functions, h_1 and h_2 :

if $h_2(n) \geq h_1(n)$

$h_2(n)$ is more informed than $h_1(n)$

Alternatively we say that $h_2(n)$ dominates $h_1(n)$.

- More informedness implies fewer expanded states.

Informed Search

Creating a heuristic function

- How to choose the heuristic function for a given problem?

Informed Search

Creating a heuristic function

- How to choose the heuristic function for a given problem?
- The heuristic function is usually chosen by means of ***problem relaxation***.

Informed Search

Creating a heuristic function

- How to choose the heuristic function for a given problem?
- The heuristic function is usually chosen by means of ***problem relaxation***.
- Problem relaxation means making the problem easier by dropping some constraints.

Informed Search

Creating a heuristic function

- How to choose the heuristic function for a given problem?
- The heuristic function is usually chosen by means of ***problem relaxation***.
- Problem relaxation means making the problem easier by dropping some constraints.
- Can also have different heuristics and always choose the best one:

$$h(n) = \max\{h_1(n), h_2(n), \dots, h_m(n)\}$$

Informed Search

Creating a heuristic function

- For the Romania map problem, the heuristic function was selected as the straight line distance between the current city and the goal.

Informed Search

Creating a heuristic function

- For the Romania map problem, the heuristic function was selected as the straight line distance between the current city and the goal.
- How is this a relaxed problem?

Informed Search

Creating a heuristic function

- For the Romania map problem, the heuristic function was selected as the straight line distance between the current city and the goal.
- How is this a relaxed problem?
 - By dropping the traveling on roads constraint.

Informed Search

Example: creating heuristic for 8-puzzle

- Consider the 8-puzzle problem.

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Consider the 8-puzzle problem.
 - It would take at least *26 moves* to solve the problem instance shown below.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Consider the 8-puzzle problem.
 - It would take at least *26 moves* to solve the problem instance shown below.
 - An admissible heuristic shouldn't overestimate that cost of 26 moves.
 - $h(\text{Start}) \leq 26$

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Consider the 8-puzzle problem.
 - It would take at least *26 moves* to solve the problem instance shown below.
 - An admissible heuristic shouldn't overestimate that cost of 26 moves.
 - $h(\text{Start}) \leq 26$
 - Generally speaking, an admissible heuristic shouldn't overestimate the cost of solving the puzzle starting from any node.
 - $h(n) \leq$ minimum number of moves to get from n to *Goal*.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:
 - $h_1(n)$ = number of tiles out of place.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:
 - $h_1(n)$ = number of tiles out of place.
 - $h_1(Start) = 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 = 8 \leq 26$.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:
 - $h_1(n)$ = number of tiles out of place.
 - $h_1(Start) = 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 = 8 \leq 26$.
 - A tile can move from loc. A to loc. B if A,B are adjacent:

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:
 - $h_1(n)$ = number of tiles out of place.
 - $h_1(Start) = 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 = 8 \leq 26$.
 - A tile can move from loc. A to loc. B if A,B are adjacent:
 - $h_2(n)$ = sum of distances of tiles from goal locations.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:
 - $h_1(n)$ = number of tiles out of place.
 - $h_1(\text{Start}) = 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 = 8 \leq 26$.
 - A tile can move from loc. A to loc. B if A,B are adjacent:
 - $h_2(n)$ = sum of distances of tiles from goal locations.
 - $h_2(\text{Start}) = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18 \leq 26$.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

Example: creating heuristic for 8-puzzle

- Description of legal moves in 8-puzzle:
A tile can move from location A to location B if A,B are adjacent and B is blank.
- Problem relaxation:
 - A tile can move from location A to location B:
 - $h_1(n)$ = number of tiles out of place.
 - $h_1(\text{Start}) = 1 + 1 + 1 + 1 + 1 + 1 + 1 + 1 = 8 \leq 26$.
 - A tile can move from loc. A to loc. B if A,B are adjacent:
 - $h_2(n)$ = sum of distances of tiles from goal locations.
 - $h_2(\text{Start}) = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18 \leq 26$.
- Notice that h_1 and h_2 are admissible, and $h_1(n) \leq h_2(n)$.

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

Informed Search

A* search - 8-puzzle search costs

Data are averaged over 100 instances for the 8-puzzle problem across various solution lengths.

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

| Optimal Solution Length (d) | Search Cost (nodes generated) | | |
|------------------------------------|-------------------------------|----------------|----------------|
| | IDS | A* using h_1 | A* using h_2 |
| 4 steps | 112 | 13 | 12 |
| 8 steps | 6384 | 39 | 25 |
| 12 steps | 3644035 | 227 | 73 |
| 16 steps | - | 1301 | 211 |
| 20 steps | - | 7276 | 676 |
| 24 steps | - | 39135 | 1641 |

Outline

Outline

Informed Search

Introduction

Greedy best-first

A*

Heuristic Functions

Requirements & Reading Material

1 Informed Search

2 Requirements & Reading Material

Requirements

What do I need from you

- When given a certain problem you should be able to:

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Build the search tree up to a given depth.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Build the search tree up to a given depth.
 - Traverse the search tree according to a given strategy.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Build the search tree up to a given depth.
 - Traverse the search tree according to a given strategy.
 - Propose a good heuristic function for the problem.
- Answer descriptive questions.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Build the search tree up to a given depth.
 - Traverse the search tree according to a given strategy.
 - Propose a good heuristic function for the problem.
 - Indicate whether a given heuristic is admissible/consistent or not.
- Answer descriptive questions.

Reading Material

Which parts of the textbook are covered

- Russell-Norvig, Chapters 3:
 - Pages 92 - 98.
 - Pages 102 - 106.