

CSE 411: Artificial Intelligence (Elective Course #6)

**400 Level, Mechatronics Engineering
2nd Term 2016/2017, Lecture #6**

Hazem Shehata

Dept. of Computer & Systems Engineering
Zagazig University

April 3rd, 2017

Credits to Dr. Mohamed El Abd for the slides

Adminstrivia

Notes

- Assignment #2:
 - Due on **Thursday**.

Course Info:

- Website: <http://hshehata.github.io/courses/zu/cse411/>
- Office hours: Sunday 11:30am - 12:30pm

Adminstrivia

Notes

- Assignment #2:
 - Due on **Thursday**.
- Assignment #3:
 - Released Today.
 - Due on **Sunday, April 16, 2017**.

Course Info:

- Website: <http://hshehata.github.io/courses/zu/cse411/>
- Office hours: Sunday 11:30am - 12:30pm

Outline

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

1 Adversarial Search

2 Requirements & Reading Material

Outline

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

1 Adversarial Search

2 Requirements & Reading Material

Adversarial search

Introduction

- Multi-agent environments:
 - Cooperative.
 - Competitive.

Adversarial search

Introduction

- Multi-agent environments:
 - Cooperative.
 - Competitive.
- In competitive environments, different agents have conflicting goals.

Adversarial search

Introduction

- Multi-agent environments:
 - Cooperative.
 - Competitive.
- In competitive environments, different agents have conflicting goals.
- This gives rise to ***adversarial search***, also known as ***games***.

Adversarial search

Introduction

- In games, other agents:
 - Want our agent to lose.
 - Introduce uncertainty, don't know what will they do.

Adversarial search

Introduction

- In games, other agents:
 - Want our agent to lose.
 - Introduce uncertainty, don't know what will they do.
- The unpredictability of other agents can introduce many possible ***contingencies***.

Adversarial search

Introduction

- In games, other agents:
 - Want our agent to lose.
 - Introduce uncertainty, don't know what will they do.
- The unpredictability of other agents can introduce many possible ***contingencies***.
- We will consider a specific type:
 - Deterministic.
 - Fully observable.
 - Two agents with alternating actions.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information	battleships, blind tictactoe	bridge, poker, scrabble nuclear war

Adversarial search

Introduction

- The main problem with games is that the search space is too huge:
 - A search tree for chess has an average branching factor of 35.
 - An average chess game lasts for 50 moves per player.
 - The average search tree has 35^{100} nodes.

Adversarial search

Introduction

- The main problem with games is that the search space is too huge:
 - A search tree for chess has an average branching factor of 35.
 - An average chess game lasts for 50 moves per player.
 - The average search tree has 35^{100} nodes.
- There is not enough time to search the whole tree and calculate the exact consequences of every move.

Adversarial search

Introduction

- The main problem with games is that the search space is too huge:
 - A search tree for chess has an average branching factor of 35.
 - An average chess game lasts for 50 moves per player.
 - The average search tree has 35^{100} nodes.
- There is not enough time to search the whole tree and calculate the exact consequences of every move.
- If the optimal decision is not feasible (due to time or memory constraints), some decision has to be taken.

Adversarial search

Introduction

We will investigate:

- What is an optimal move and how to find it.

Adversarial search

Introduction

We will investigate:

- What is an optimal move and how to find it.
- Pruning the search tree.

Adversarial search

Introduction

We will investigate:

- What is an optimal move and how to find it.
- Pruning the search tree.
- How to choose a good move when time is limited.

Adversarial search

Introduction

We will investigate:

- What is an optimal move and how to find it.
- Pruning the search tree.
- How to choose a good move when time is limited.
- How to define a heuristic evaluation function to approximate the utility of a state.

Adversarial search

Minimax Algorithm

- First thing to consider is how to make an optimal move.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- First thing to consider is how to make an optimal move.
- Certain assumptions exist:
 - A two player game.
 - The play is sequential (agents taking turns).
 - The opponent is playing perfectly.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- First thing to consider is how to make an optimal move.
- Certain assumptions exist:
 - A two player game.
 - The play is sequential (agents taking turns).
 - The opponent is playing perfectly.
- The two players are referred to as **MAX** and **MIN**:
 - **MAX** wants to maximize the utility.
 - **MIN** wants to minimize the utility.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.
 - $\text{PLAYER}(s)$: which player has the move in a state.

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.
 - $\text{PLAYER}(s)$: which player has the move in a state.
 - $\text{ACTIONS}(s)$: set of legal moves in a state.

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.
 - $\text{PLAYER}(s)$: which player has the move in a state.
 - $\text{ACTIONS}(s)$: set of legal moves in a state.
 - $\text{RESULT}(s, a)$: **transition model**; result of a move.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.
 - $\text{PLAYER}(s)$: which player has the move in a state.
 - $\text{ACTIONS}(s)$: set of legal moves in a state.
 - $\text{RESULT}(s, a)$: **transition model**; result of a move.
 - $\text{TERMINAL-TEST}(s)$: **terminal test**; true if and only if the game is over. States where game ends are called **terminal states**.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.
 - $\text{PLAYER}(s)$: which player has the move in a state.
 - $\text{ACTIONS}(s)$: set of legal moves in a state.
 - $\text{RESULT}(s, a)$: **transition model**; result of a move.
 - $\text{TERMINAL-TEST}(s)$: **terminal test**; true if and only if the game is over. States where game ends are called **terminal states**.
 - $\text{UTILITY}(s)$: **utility function** (a.k.a., payoff function or objective function); final numeric value for a game that ends in a terminal state s .

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- A **game** can be formally defined as a kind of search problem with the following elements :
 - S_0 : **initial state**; how game is set up at start.
 - $\text{PLAYER}(s)$: which player has the move in a state.
 - $\text{ACTIONS}(s)$: set of legal moves in a state.
 - $\text{RESULT}(s, a)$: **transition model**; result of a move.
 - $\text{TERMINAL-TEST}(s)$: **terminal test**; true if and only if the game is over. States where game ends are called **terminal states**.
 - $\text{UTILITY}(s)$: **utility function** (a.k.a., payoff function or objective function); final numeric value for a game that ends in a terminal state s .
- S_0 , ACTIONS , and RESULT define the **game tree**.

Outline

Adversarial Search

Introduction

Minimax Algorithm


α - β Pruning

Requirements & Reading Material

Adversarial search

Example: game tree for tic-tac-toe

A partial game tree for the game of tic-tac-toe:



Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- **MAX** is the first player to move, starting at the root node.

Adversarial search

Minimax Algorithm

- **MAX** is the first player to move, starting at the root node.
- The utility values at the leaf nodes are from **MAX's** point of view.

Adversarial search

Minimax Algorithm

- **MAX** is the first player to move, starting at the root node.
- The utility values at the leaf nodes are from **MAX's** point of view.
- The optimal strategy is found by examining the ***minimax value*** of each node.

Adversarial search

Minimax Algorithm

- In a normal search tree, a solution is to find a path from initial state to goal state (*i.e.*, a winning state for **MAX**).

Adversarial search

Minimax Algorithm

- In a normal search tree, a solution is to find a path from initial state to goal state (*i.e.*, a winning state for **MAX**).
- However, since **MIN** affects the search, a contingent strategy should be found.

Adversarial search

Minimax Algorithm

- In a normal search tree, a solution is to find a path from initial state to goal state (*i.e.*, a winning state for **MAX**).
- However, since **MIN** affects the search, a contingent strategy should be found.
- The strategy defines:
 - **MAX's** move at the initial state.
 - **MAX's** move at all states generated by all possible moves by **MIN**.
 - ... and so on.

Adversarial search

Minimax Algorithm

- In a normal search tree, a solution is to find a path from initial state to goal state (*i.e.*, a winning state for **MAX**).
- However, since **MIN** affects the search, a contingent strategy should be found.
- The strategy defines:
 - **MAX's** move at the initial state.
 - **MAX's** move at all states generated by all possible moves by **MIN**.
 - ... and so on.
- This means that **MAX** needs a winning strategy independent of what **MIN** does.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- **MAX** will choose a move that maximizes its utility value.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- **MAX** will choose a move that maximizes its utility value.
- **MIN** will choose a move that minimizes the utility value of **MAX**.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- **MAX** will choose a move that maximizes its utility value.
- **MIN** will choose a move that minimizes the utility value of **MAX**.
- This results in **MAX** (at the root) choosing its best move given available information.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- **MAX** will choose a move that maximizes its utility value.
- **MIN** will choose a move that minimizes the utility value of **MAX**.
- This results in **MAX** (at the root) choosing its best move given available information.
- Available information is via look ahead (the search down the tree).

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Minimax Algorithm

- The minimax value of a node is defined as follows:

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Outline

Adversarial Search

Introduction

Minimax Algorithm

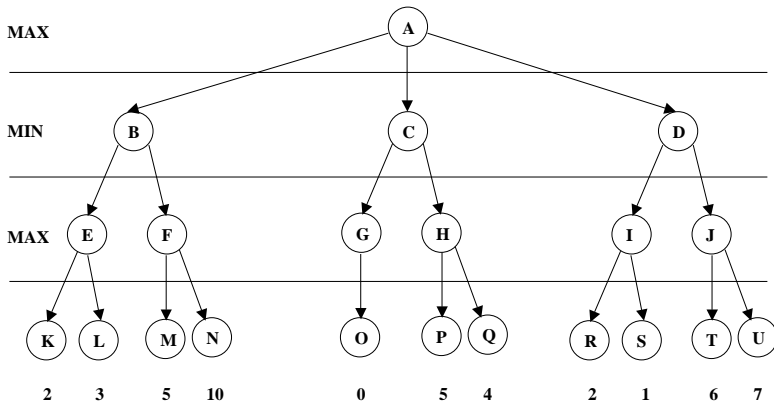
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax

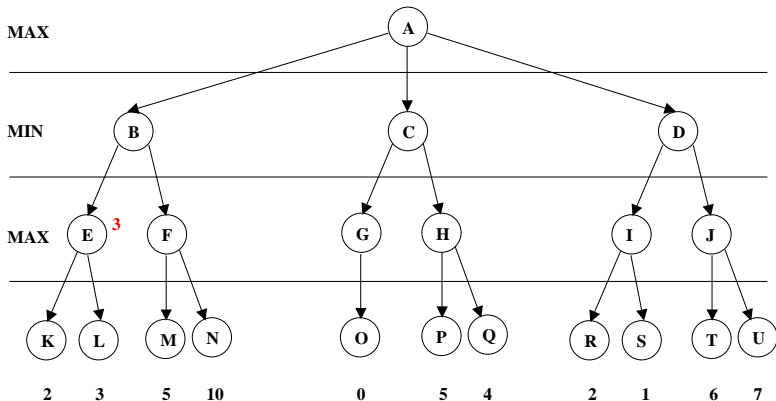
Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

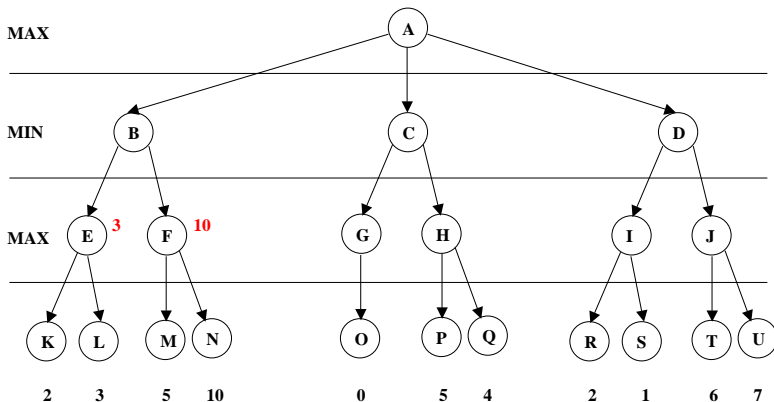
Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax


Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

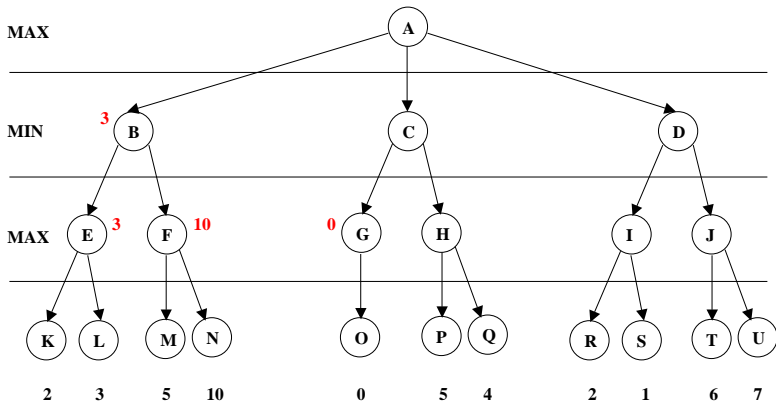
Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

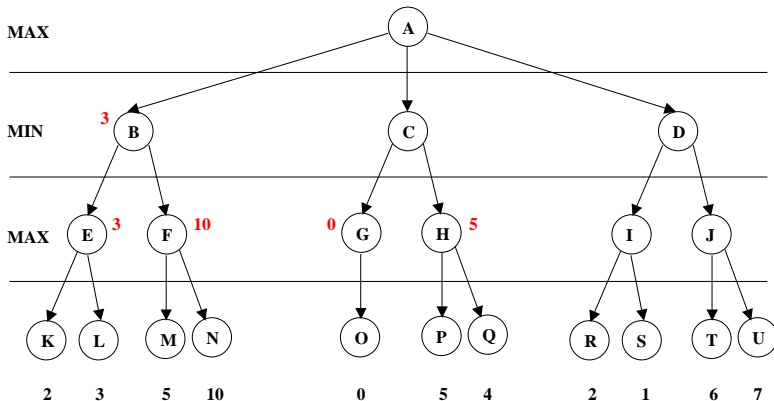
Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

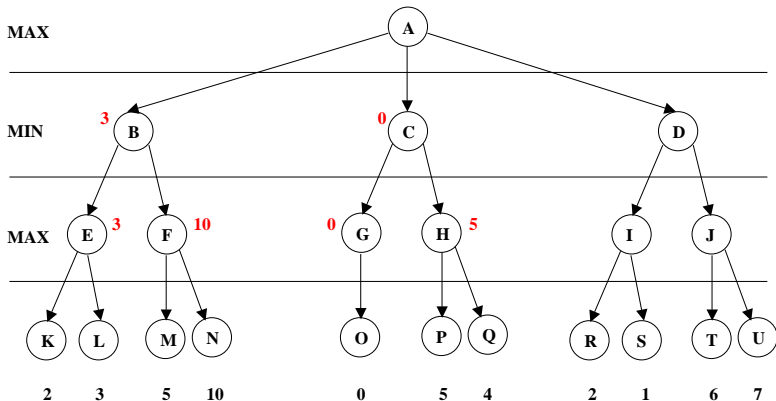
Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

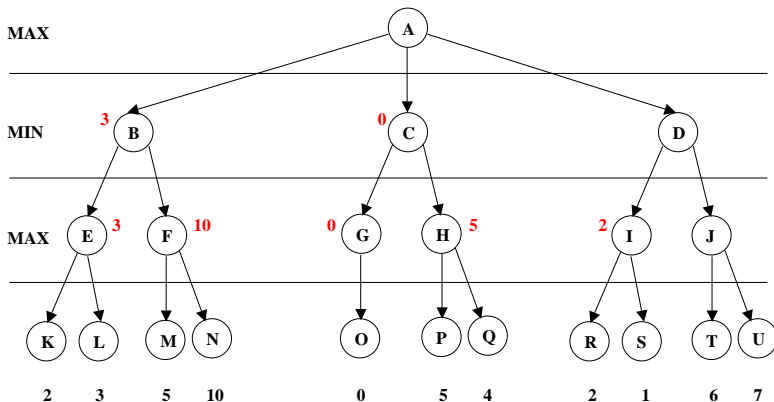
Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

Calculate the minimax value for the node A.



Outline

Adversarial
Search

Introduction

Minimax Algorithm

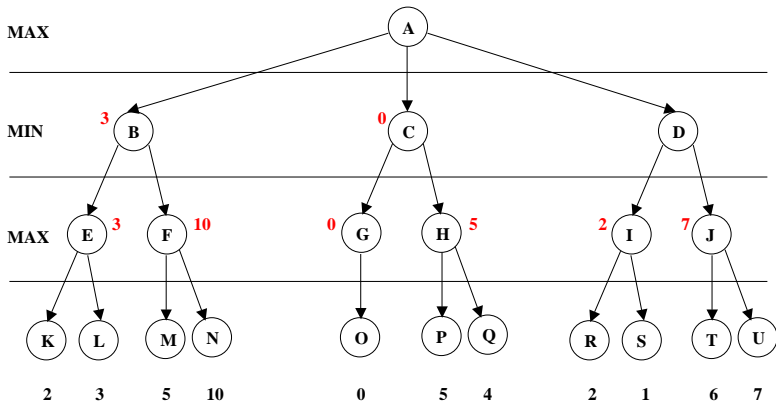
α - β Pruning

Requirements
& Reading
Material

Adversarial search

Example: applying minimax

Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

Calculate the minimax value for the node A.

Outline

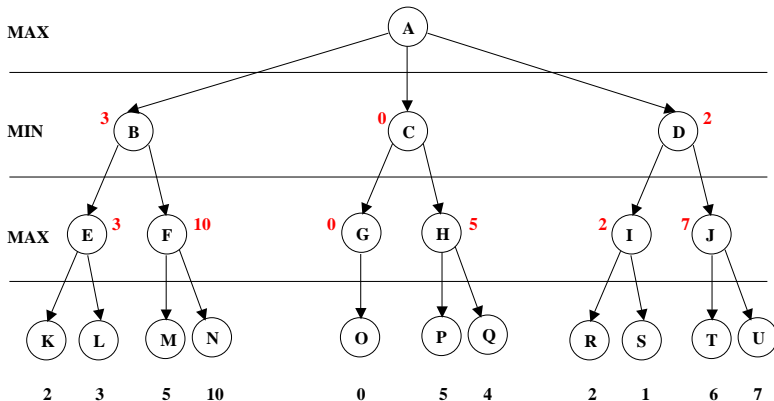
Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

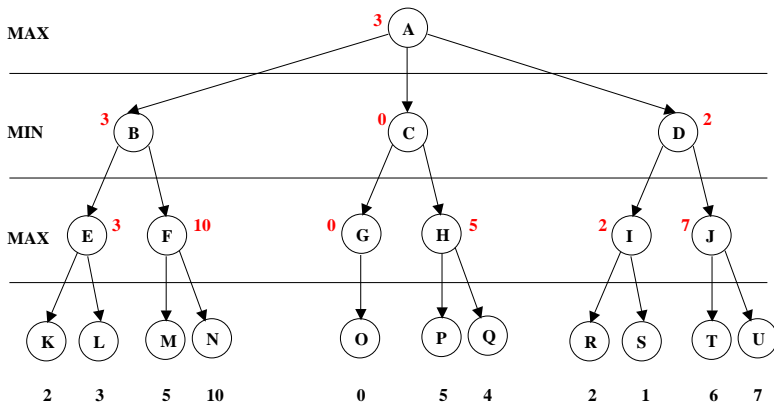
Requirements & Reading Material



Adversarial search

Example: applying minimax


Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax


Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax


Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax

Calculate the minimax value for the node A.



Adversarial search

Example: applying minimax to NIM

- Consider the game of NIM:
 - 7 matches placed in a pile.
 - Each player divides a pile of matches into 2 non-empty piles with a different number of matches.
 - The player who cannot make a move is the loser.

Adversarial search

Example: applying minimax to NIM

- Consider the game of NIM:
 - 7 matches placed in a pile.
 - Each player divides a pile of matches into 2 non-empty piles with a different number of matches.
 - The player who cannot make a move is the loser.
- The utility is +1 (**MAX** wins) or -1 (**MAX** loses).

Adversarial search

Example: applying minimax to NIM

- Consider the game of NIM:
 - 7 matches placed in a pile.
 - Each player divides a pile of matches into 2 non-empty piles with a different number of matches.
 - The player who cannot make a move is the loser.
- The utility is +1 (**MAX** wins) or -1 (**MAX** loses).
- The value at each node represents the best value of the best terminal state the current player can hope to achieve.

Outline

Adversarial Search

Introduction

Minimax Algorithm

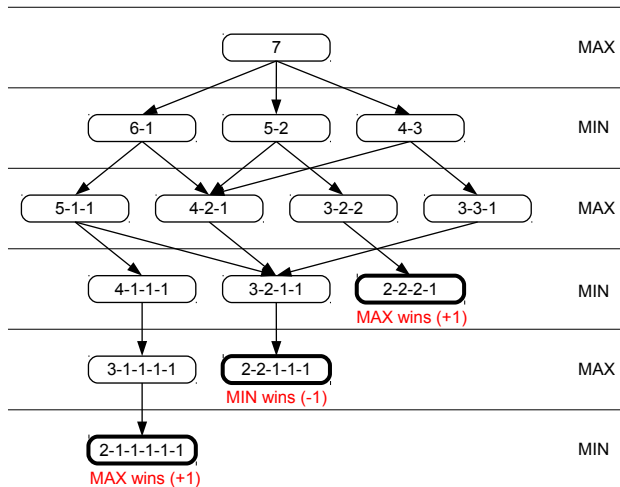
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

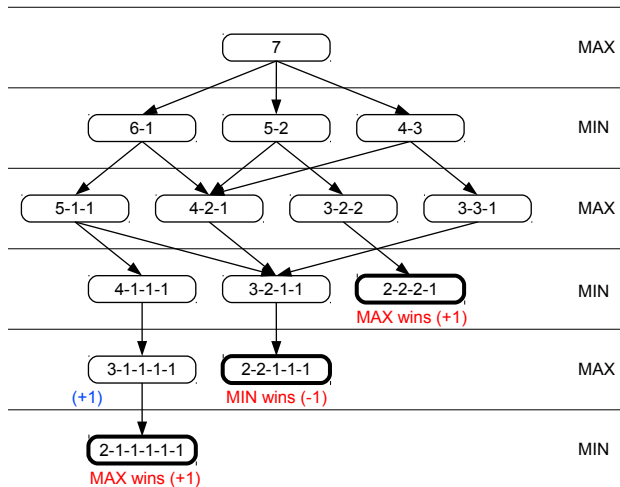
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

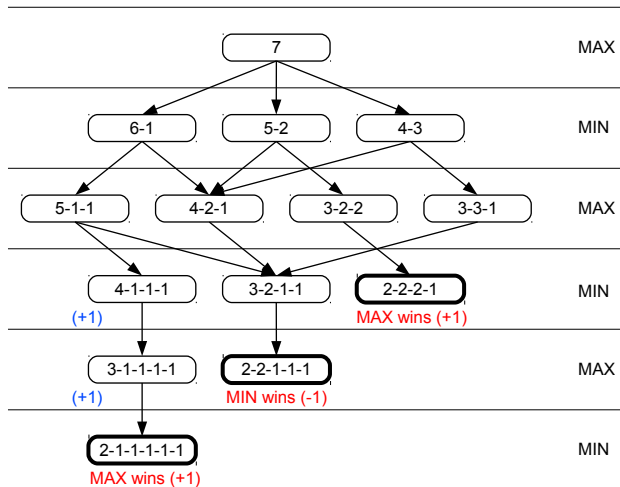
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

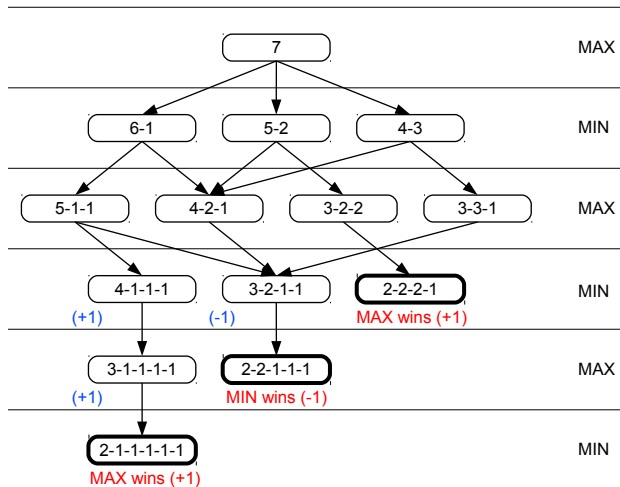
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

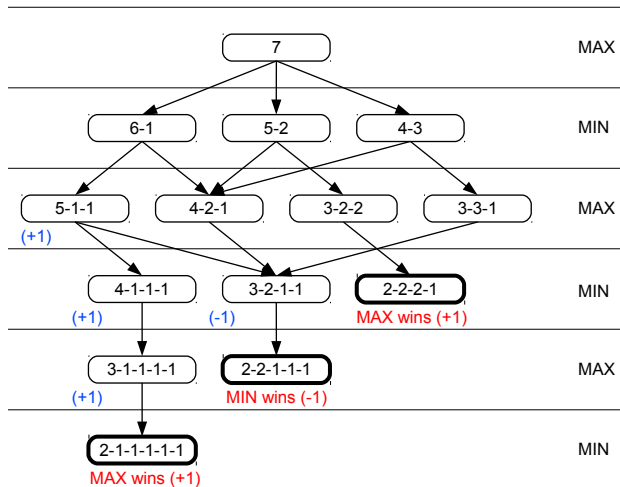
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

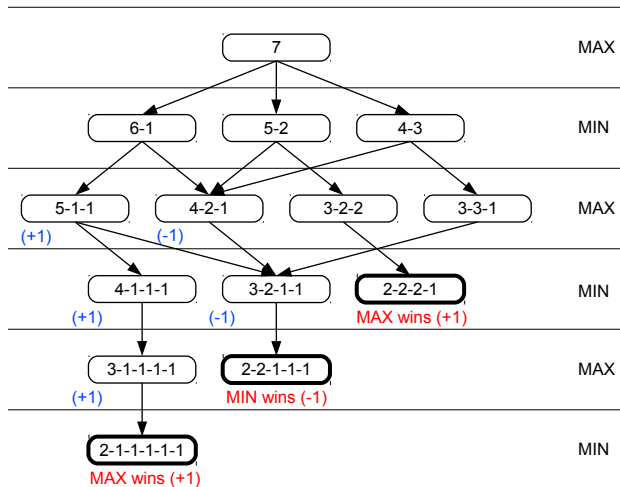
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

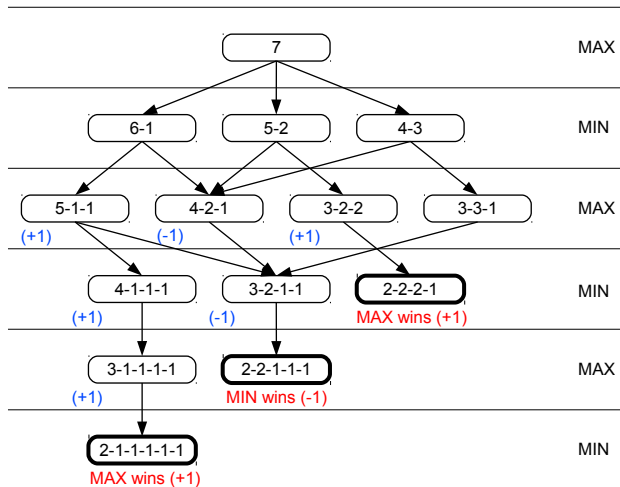
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

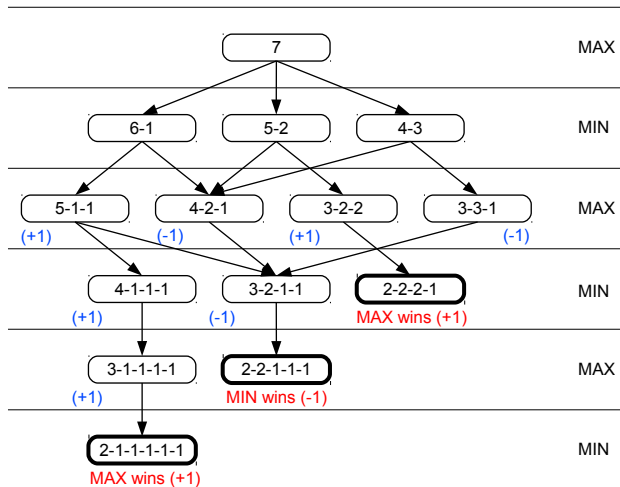
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

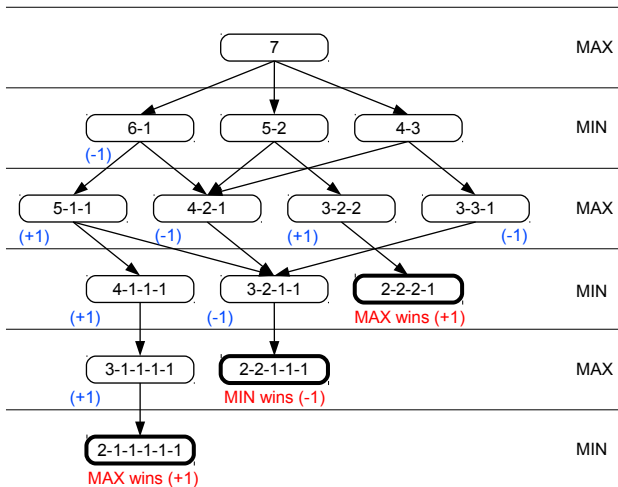
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.

Outline

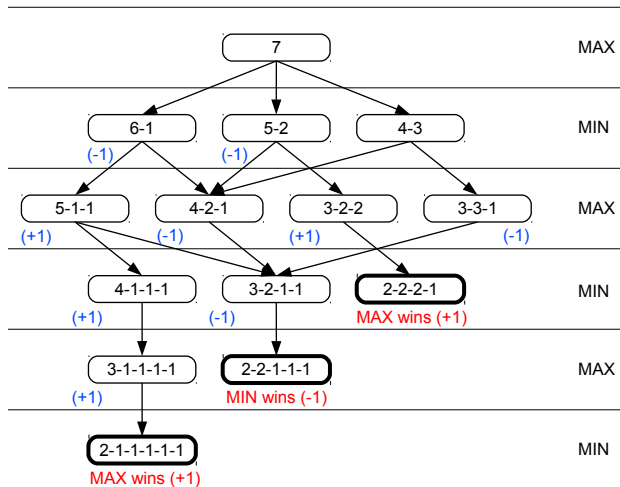
Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

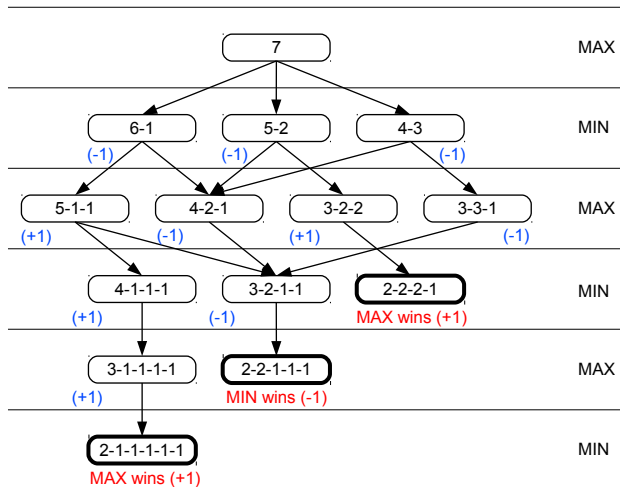
Requirements & Reading Material



Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

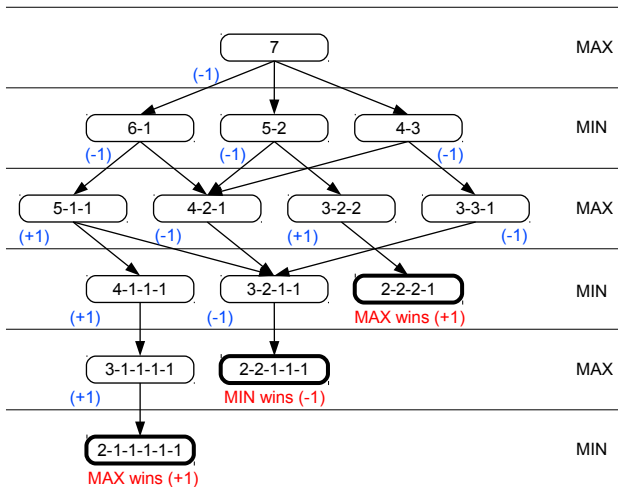
α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying minimax to NIM

Calculate the minimax value for the root node.



Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

The minimax search algorithm

function MINIMAX-SEARCH(s) **returns** an action

$v \leftarrow \text{MAX-VALUE}(s)$

return the **action** in ACTIONS(s) with value v

Adversarial search

The minimax search algorithm

function MINIMAX-SEARCH(s) **returns** an action

$v \leftarrow \text{MAX-VALUE}(s)$

return the **action** in $\text{ACTIONS}(s)$ with value v

function MAX-VALUE(s) **returns** a utility value

if $\text{TERMINAL-TEST}(s)$ **then return** $\text{UTILITY}(s)$

$v \leftarrow -\infty$

for each a **in** $\text{ACTIONS}(s)$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$

return v

Adversarial search

The minimax search algorithm

function MINIMAX-SEARCH(s) **returns** an action

$v \leftarrow \text{MAX-VALUE}(s)$

return the **action** in ACTIONS(s) with value v

function MAX-VALUE(s) **returns** a utility value

if TERMINAL-TEST(s) **then return** UTILITY(s)

$v \leftarrow -\infty$

for each a **in** ACTIONS(s) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$

return v

function MIN-VALUE(s) **returns** a utility value

if TERMINAL-TEST(s) **then return** UTILITY(s)

$v \leftarrow +\infty$

for each a **in** ACTIONS(s) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$

return v

Adversarial search

Properties of the minimax search algorithm

- Complete if the tree is finite.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Properties of the minimax search algorithm

- Complete if the tree is finite.
- Optimal against an optimal opponent.

Adversarial search

Properties of the minimax search algorithm

- Complete if the tree is finite.
- Optimal against an optimal opponent.
 - No other strategy could do any better in this case.

Adversarial search

Properties of the minimax search algorithm

- Complete if the tree is finite.
- Optimal against an optimal opponent.
 - No other strategy could do any better in this case.
 - For non-optimal opponents, outcome cannot be worse but there could be another strategy with better outcome!

Adversarial search

Properties of the minimax search algorithm

- Complete if the tree is finite.
- Optimal against an optimal opponent.
 - No other strategy could do any better in this case.
 - For non-optimal opponents, outcome cannot be worse but there could be another strategy with better outcome!
- Time complexity = $O(b^m)$.

Adversarial search

Properties of the minimax search algorithm

- Complete if the tree is finite.
- Optimal against an optimal opponent.
 - No other strategy could do any better in this case.
 - For non-optimal opponents, outcome cannot be worse but there could be another strategy with better outcome!
- Time complexity = $O(b^m)$.
- Space complexity = $O(bm)$.

Adversarial search

α - β Pruning

- Game trees are too huge to be solved to optimality.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

α - β Pruning

- Game trees are too huge to be solved to optimality.
- Minimax has to generate the entire tree, compute the utility values at the leaf nodes, and then propagate these values up the tree.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

α - β Pruning

- Game trees are too huge to be solved to optimality.
- Minimax has to generate the entire tree, compute the utility values at the leaf nodes, and then propagate these values up the tree.
- We want a way to do less work by removing unnecessary branches thus exploring less nodes.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

α - β Pruning

- Game trees are too huge to be solved to optimality.
- Minimax has to generate the entire tree, compute the utility values at the leaf nodes, and then propagate these values up the tree.
- We want a way to do less work by removing unnecessary branches thus exploring less nodes.
- However, we want to guarantee the same decision as Minimax.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

α - β Pruning

- Game trees are too huge to be solved to optimality.
- Minimax has to generate the entire tree, compute the utility values at the leaf nodes, and then propagate these values up the tree.
- We want a way to do less work by removing unnecessary branches thus exploring less nodes.
- However, we want to guarantee the same decision as Minimax.
- We use the **alpha-beta (α - β) pruning** technique.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

Example: illustrating α - β pruning

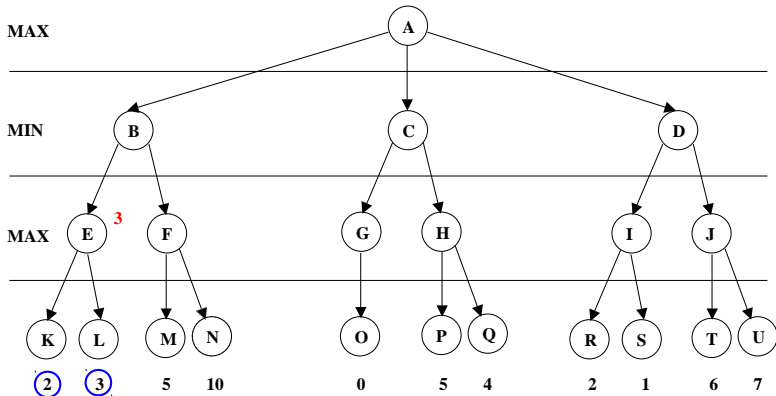
Show a scenario where α - β pruning becomes applicable.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search

Example: illustrating α - β pruning

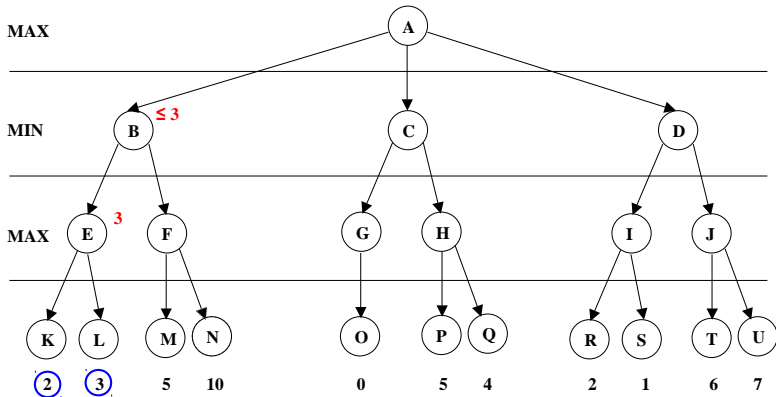
Show a scenario where α - β pruning becomes applicable.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search

Example: illustrating α - β pruning

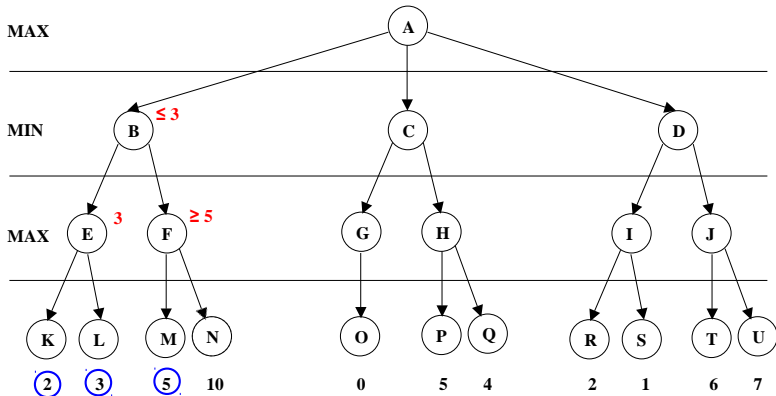
Show a scenario where α - β pruning becomes applicable.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search

Example: illustrating α - β pruning

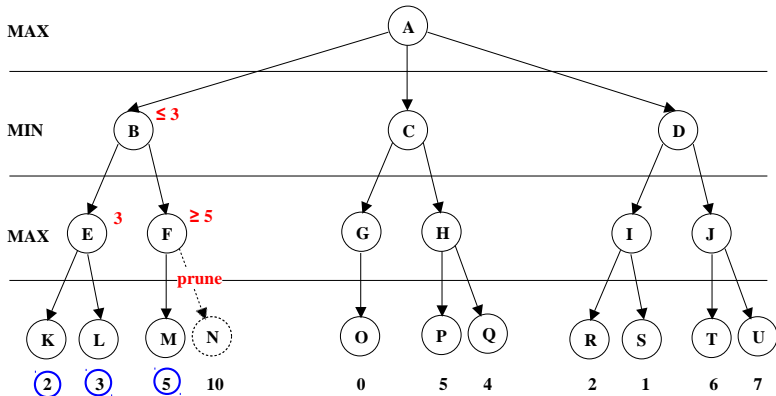
Show a scenario where α - β pruning becomes applicable.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search

Outline


Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material

α -cutoff

- α is the best (maximum) value **MAX** is assured so far off the current path.
- If β is worse than α (from **MAX's** point of view), **MAX** will avoid it, dashed branches will be pruned.
- Happens when $\alpha \geq \beta$, referred to as α -cutoff.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search


Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.

α : Best already explored option
along path to root for MAX
 β : Best already explored option
along path to root for MIN

$v = -\infty$
 $\alpha = -\infty$
 $\beta = +\infty$

v : Best already explored option
among children




Apply α - β pruning to calculate minimax value for node A.

Adversarial Search

- Introduction
- Minimax Algorithm
- α - β Pruning


Requirements & Reading Material



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning


Requirements & Reading Material



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning


Requirements & Reading Material



Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Outline

Adversarial Search


Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Outline

Adversarial Search


Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material



Adversarial search

Example: applying α - β pruning


Apply α - β pruning to calculate minimax value for node A.

Outline

Adversarial Search

Introduction
Minimax Algorithm
 α - β Pruning


Requirements & Reading Material



Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Outline

Adversarial Search


Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Outline

Adversarial Search


Introduction
Minimax Algorithm
 α - β Pruning

Requirements & Reading Material

Adversarial search

Example: applying α - β pruning

Apply α - β pruning to calculate minimax value for node A.



Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

α - β Pruning

- In the previous example, α - β pruning explored 14 nodes instead of 21 nodes explored by Minimax.

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

α - β Pruning

- In the previous example, α - β pruning explored 14 nodes instead of 21 nodes explored by Minimax.
- α - β pruning also reached the same decision reached by Minimax.

Adversarial search

The α - β search Algorithm

function ALPHA-BETA-SEARCH(s) **returns** an action

$v \leftarrow \text{MAX-VALUE}(s, -\infty, +\infty)$

return the **action** in ACTIONS(s) with value v

Adversarial search

The α - β search Algorithm

function ALPHA-BETA-SEARCH(s) **returns** an action

$v \leftarrow \text{MAX-VALUE}(s, -\infty, +\infty)$

return the **action** in ACTIONS(s) with value v

function MAX-VALUE(s, α, β) **returns** a utility value

if TERMINAL-TEST(s) **then return** UTILITY(s)

$v \leftarrow -\infty$

for each a **in** ACTIONS(s) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if $\alpha \geq \beta$ **then return** α /* β -cutoff */

return v

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Adversarial search

The α - β search Algorithm

function ALPHA-BETA-SEARCH(s) **returns** an action

$v \leftarrow \text{MAX-VALUE}(s, -\infty, +\infty)$

return the **action** in ACTIONS(s) with value v

function MAX-VALUE(s, α, β) **returns** a utility value

if TERMINAL-TEST(s) **then return** UTILITY(s)

$v \leftarrow -\infty$

for each a **in** ACTIONS(s) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if $\alpha \geq \beta$ **then return** α */* β -cutoff */*

return v

function MIN-VALUE(s, α, β) **returns** a utility value

if TERMINAL-TEST(s) **then return** UTILITY(s)

$v \leftarrow +\infty$

for each a **in** ACTIONS(s) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

$\beta \leftarrow \text{MIN}(\beta, v)$

if $\alpha \geq \beta$ **then return** β */* α -cutoff */*

return v

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

Outline

Outline

Adversarial Search

Introduction

Minimax Algorithm

α - β Pruning

Requirements & Reading Material

1 Adversarial Search

2 Requirements & Reading Material

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Formulate the game problem.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Formulate the game problem.
 - Build the game tree up to a given depth.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Formulate the game problem.
 - Build the game tree up to a given depth.
 - Apply the minimax search algorithm.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Formulate the game problem.
 - Build the game tree up to a given depth.
 - Apply the minimax search algorithm.
 - Apply the $\alpha - \beta$ pruning algorithm.

Requirements

What do I need from you

- When given a certain problem you should be able to:
 - Formulate the game problem.
 - Build the game tree up to a given depth.
 - Apply the minimax search algorithm.
 - Apply the $\alpha - \beta$ pruning algorithm.
- Answer descriptive questions.

Reading Material

Which parts of the textbook are covered

- Russell-Norvig, Chapters 5:
 - Pages 161 - 170.